

OpenShift 4

For Developers

Wanja Pernath
wpernath@redhat.com

EMEA Partner Enablement Manager, OpenShift & MW

Name: Wanja Pernath

Email: wpernath@redhat.com

Base: Germany (very close to the Alps)

Role: EMEA Technical Partner Development

Manager - OpenShift and MW

Experience: Years of Consulting, Training,

PreSales at Red Hat and before

Self introduction



OpenShift Container Platform

Advanced Cluster Management

Multi-cluster Management

Discovery : Policy : Compliance : Configuration : Workloads

OpenShift Container Platform

Manage Workloads

Platform Services

Service Mesh : Serverless
Builds : CI/CD Pipelines
Full Stack Logging
Chargeback

Build Cloud-Native Apps

Application Services

Databases : Languages
Runtimes : Integration
Business Automation
100+ ISV Services

Developer Productivity

Developer Services

Developer CLI : VS Code extensions : IDE Plugins
Code Ready Workspaces
CodeReady Containers

Cluster Services

Automated Ops : Over-The-Air Updates : Monitoring : Registry : Networking : Router : KubeVirt : OLM : Helm

Kubernetes

Red Hat Enterprise Linux & RHEL CoreOS



Physical



Virtual



Private cloud



Public cloud



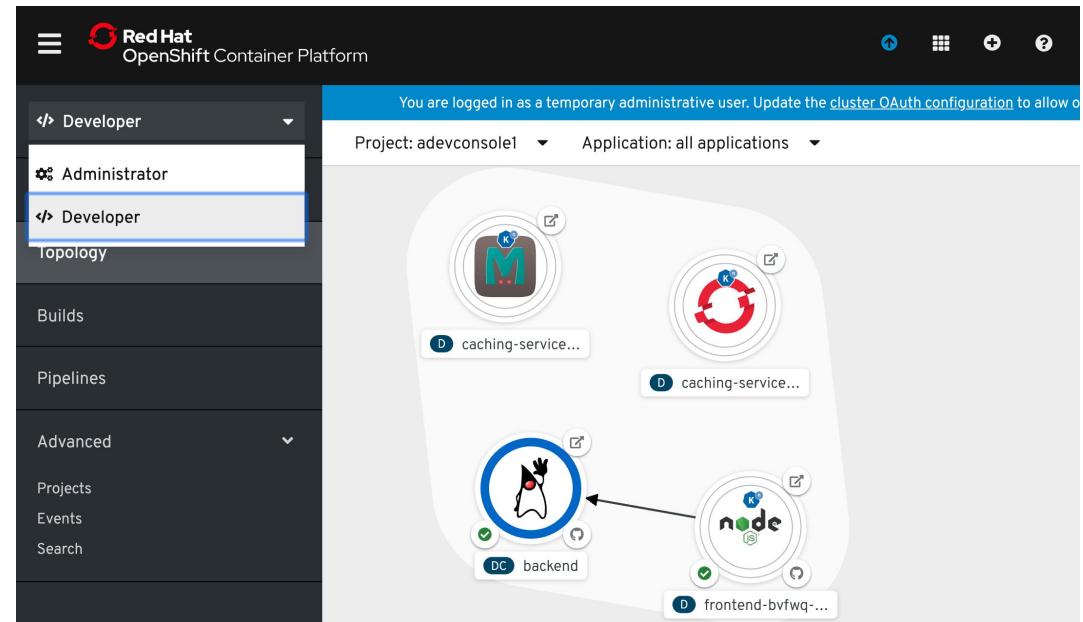
Managed cloud
(Azure, AWS, IBM, Red Hat)

Web Console

OpenShift Developer Console

An developer-specific perspective in the OpenShift UI that will sit beside the admin console and focus on developer use cases.

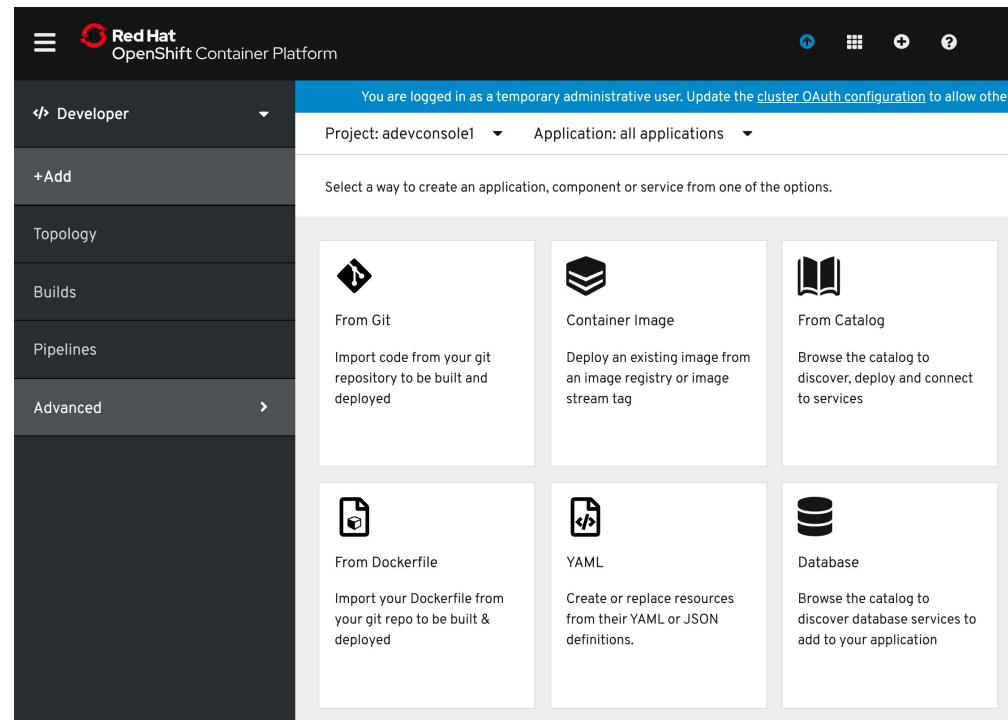
All OpenShift developer tool UIs will be surfaced here.



Developer Console: Create Applications

Key Features

- Import source from Git
- View existing container image
- Edit YAML definition
- Build from Dockerfile
- Explore services catalog
- Deploy database from catalog



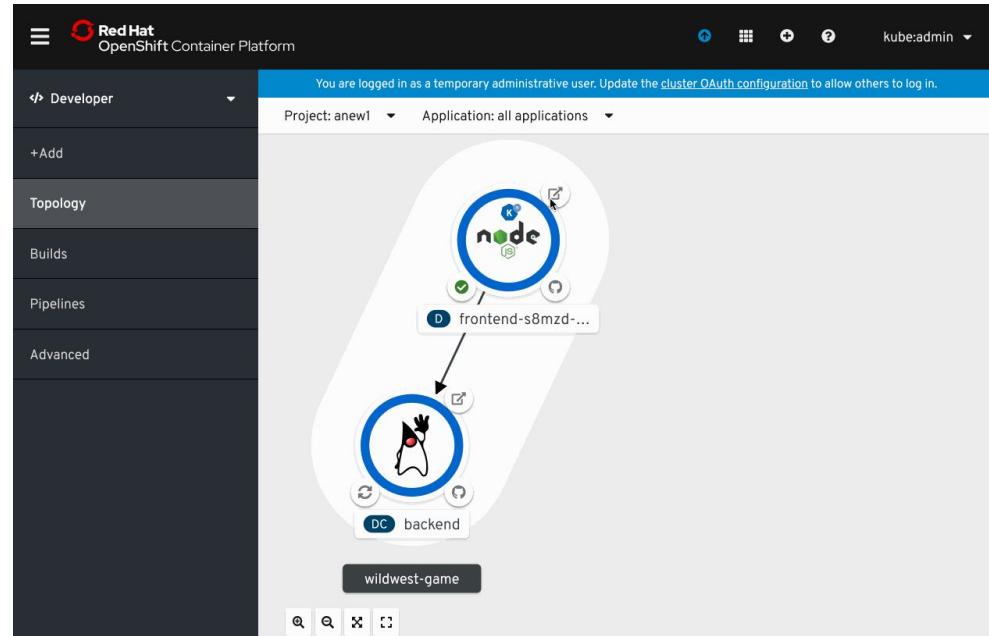
The screenshot shows the Red Hat OpenShift Container Platform Developer Console. The top navigation bar includes the Red Hat logo and the text "Red Hat OpenShift Container Platform". The top right corner has icons for user profile, grid, add, and help. The main header bar has a "Developer" dropdown, a "Project: adevconsole1" dropdown, and an "Application: all applications" dropdown. A blue banner at the top right says "You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow other". The left sidebar has a "Developer" section with "Add" (highlighted in blue), "Topology", "Builds", "Pipelines", and "Advanced". The main content area is titled "Select a way to create an application, component or service from one of the options." It contains six cards arranged in a 2x3 grid:

- From Git**: Import code from your git repository to be built and deployed.
- Container Image**: Deploy an existing image from an image registry or image stream tag.
- From Catalog**: Browse the catalog to discover, deploy and connect to services.
- From Dockerfile**: Import your Dockerfile from your git repo to be built & deployed.
- YAML**: Create or replace resources from their YAML or JSON definitions.
- Database**: Browse the catalog to discover database services to add to your application.

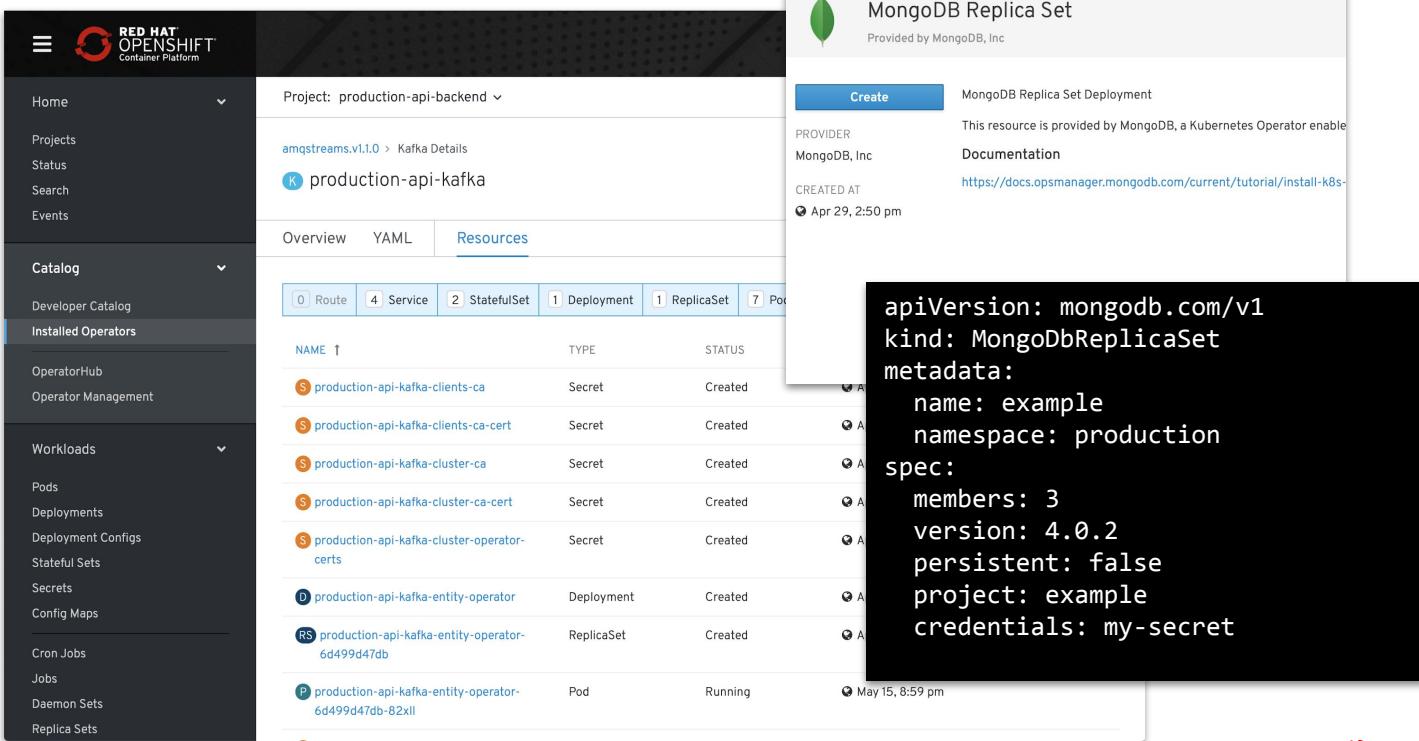
Developer Console: Application Topology

Key Features

- View structure and status of app components
- Drill into specific workloads
- Quickly navigate to pod logs
- Manually scale
- Pod donut!
- Access route/URL
- Linked build and source



Self-service for developers

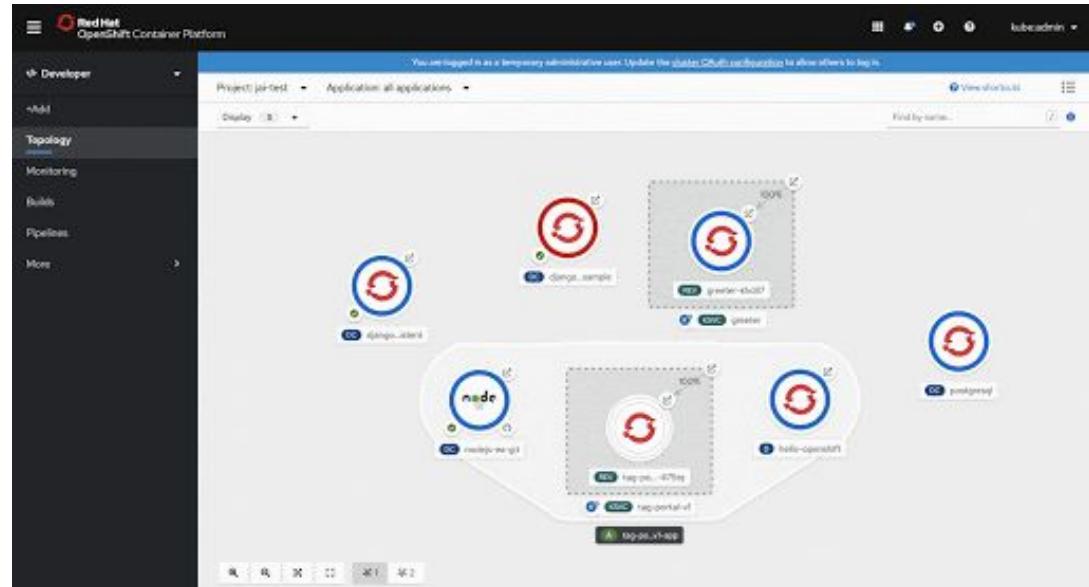


The screenshot shows the Red Hat OpenShift v4 web interface. The left sidebar includes sections for Home, Projects, Status, Search, Events, Catalog (Developer Catalog, Installed Operators selected), OperatorHub, and Operator Management. The main content area shows a project named 'production-api-backend' with a Kafka Details view. A tooltip for the 'MongoDB Replica Set' operator is open, showing its provider (MongoDB, Inc), creation date (Apr 29, 2:50 pm), and a link to its documentation. Below the tooltip, a snippet of the operator's configuration YAML is displayed:

```
apiVersion: mongodb.com/v1
kind: MongoDBReplicaSet
metadata:
  name: example
  namespace: production
spec:
  members: 3
  version: 4.0.2
  persistent: false
  project: example
  credentials: my-secret
```

Understanding Complex Apps

- Allow developers to easily find and focus on interested components
- Matched items are highlighted and non-matches are dimmed
- Indicates matches off of visible area



Understanding Complex Apps grouping, collapse

Operator-managed and Helm-based services

- Groups all resources to easily understand what is being managed by which CR/operator
- Add simple info in resources side panel, for custom resource or general resource parent

Collapse / Expand

- Manage large multi-application projects by expanding/collapsing certain grouping types
- Reduce Application, Helm, Knative and Operator to simple summary tiles

OpenShift Container Platform

Advanced Cluster Management

Multi-cluster Management

Discovery : Policy : Compliance : Configuration : Workloads

OpenShift Container Platform

Manage Workloads

Platform Services

Service Mesh : Serverless
Builds : CI/CD Pipelines
Full Stack Logging
Chargeback

Build Cloud-Native Apps

Application Services

Databases : Languages
Runtimes : Integration
Business Automation
100+ ISV Services

Developer Productivity

Developer Services

Developer CLI : VS Code extensions : IDE Plugins
Code Ready Workspaces
CodeReady Containers

OpenShift Kubernetes Engine

Cluster Services

Automated Ops : Over-The-Air Updates : Monitoring : Registry : Networking : Router : KubeVirt : OLM : Helm

Kubernetes

Red Hat Enterprise Linux & RHEL CoreOS



Physical



Virtual



Private cloud



Public cloud



Managed cloud
(Azure, AWS, IBM, Red Hat)



Desktop Tools, IDEs & Plugins

CodeReady Workspaces 2.x

- New onboarding experience - updated catalog and simple workspace creation options
- Quarkus support - includes stack and plugin
- Updated languages - .NET 3.1, Camel K, Java 11 & Gradle
- Air Gap improvements
- Many editor enhancements

Getting Started with CodeReady Workspaces

Select a Sample

Use a sample to create your first workspace.

 JBoss Java stack with JBoss 7.2, OpenShift 3.6 and Maven 3.5	 Thortail REST HTTP	 Red Hat Fuse	 Java Maven	 Apache Camel K	 Java Gradle
 Quarkus Tools	 Java Spring Boot	 Java Vert.x	 Node.js Express Web Application	 Node.js ConfigMap	 Node.js MongoDB Web Application
 C++	 .NET	 Go	 CakePHP Example	 PHP Simple	 Python

CREATE A NEW

Temporary Storage 1000MB 30GB 18 Items

Search

 QUARKUS

Quarkus Tools

Quarkus Tools with GraalVM and Maven 3.6.0

>qui

Quarkus: Add extensions to current project

Quarkus: Debug current Quarkus project

Quarkus: Generate a Quarkus project

Quarkus: Welcome

Developer-focused CLI for rapid development iterations on OpenShift

Simplifies building of microservices applications on OpenShift.

odo - OpenShift's Dev-Focused CLI

```
$ odo create wildfly backend
Component 'backend' was created.

$ odo push
Pushing changes to component: backend

$ odo create php frontend
Component 'frontend' was created.
To push source code to the component run 'odo push'

$ odo push
Pushing changes to component: frontend

$ odo url create
frontend - http://frontend-myapp.192.168.99.100.nip.io

$ odo watch
Waiting for something to change in /dev/frontend
```

CodeReady Containers

Environment to simplify local direct-to-OpenShift application development.

Available for:

- Linux (KVM)
- Windows (Virtualbox)
- MacOS (Virtualbox)

```
$ crc setup  
Prepare your machine for running OpenShift
```

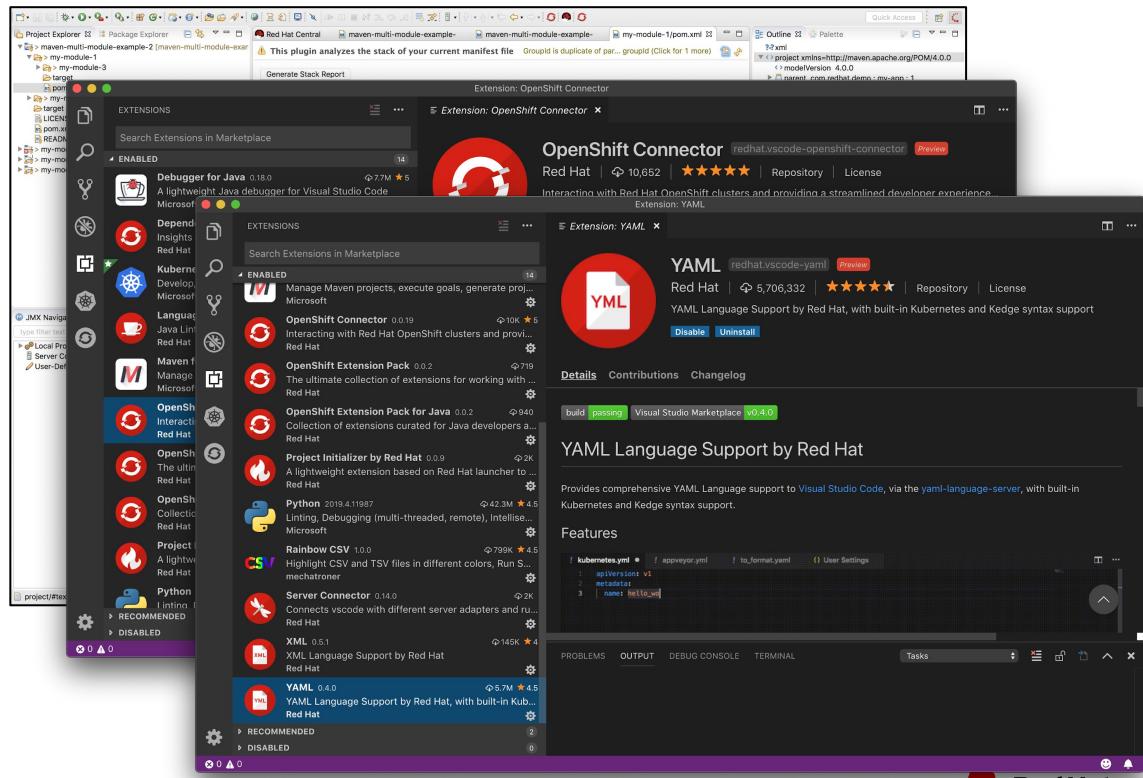
```
$ crc start -b  
crc-hyperkit-4.2.0.crcbundle  
Start with the Hyperkit 4.2 bundle
```

```
$ crc status  
Get the status of the cluster
```

IDE Plugins for OpenShift

We have a suite of plugins to simplify working with OpenShift and YAML, as well as our Middleware products such as Java, Camel, Wildfly, and much more. These extensions are available for:

- Azure DevOps
- VS Code
- JetBrains IDEs (e.g. IntelliJ)
- CodeReady Workspaces
- CodeReady Studio
- Eclipse IDE/JBoss Tools



OpenShift Container Platform

Advanced Cluster Management

Multi-cluster Management

Discovery : Policy : Compliance : Configuration : Workloads

OpenShift Container Platform

Manage Workloads

Platform Services

Service Mesh : Serverless
Builds : CI/CD Pipelines
Full Stack Logging
Chargeback

Build Cloud-Native Apps

Application Services

Databases : Languages
Runtimes : Integration
Business Automation
100+ ISV Services

Developer Productivity

Developer Services

Developer CLI : VS Code extensions : IDE Plugins
Code Ready Workspaces
CodeReady Containers

Cluster Services

Automated Ops : Over-The-Air Updates : Monitoring : Registry : Networking : Router : KubeVirt : OLM : Helm

Kubernetes

Red Hat Enterprise Linux & RHEL CoreOS



Physical



Virtual



Private cloud



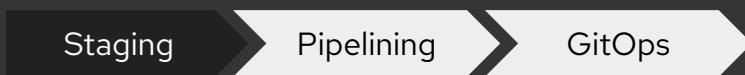
Public cloud



Managed cloud
(Azure, AWS, IBM, Red Hat)



From Code to Image

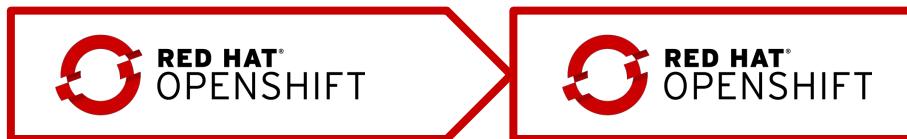


Red Hat OpenShift

Adds consistency and portability to your cloud journey

DEVELOPMENT

PRODUCTION



Use OpenShift across any environment



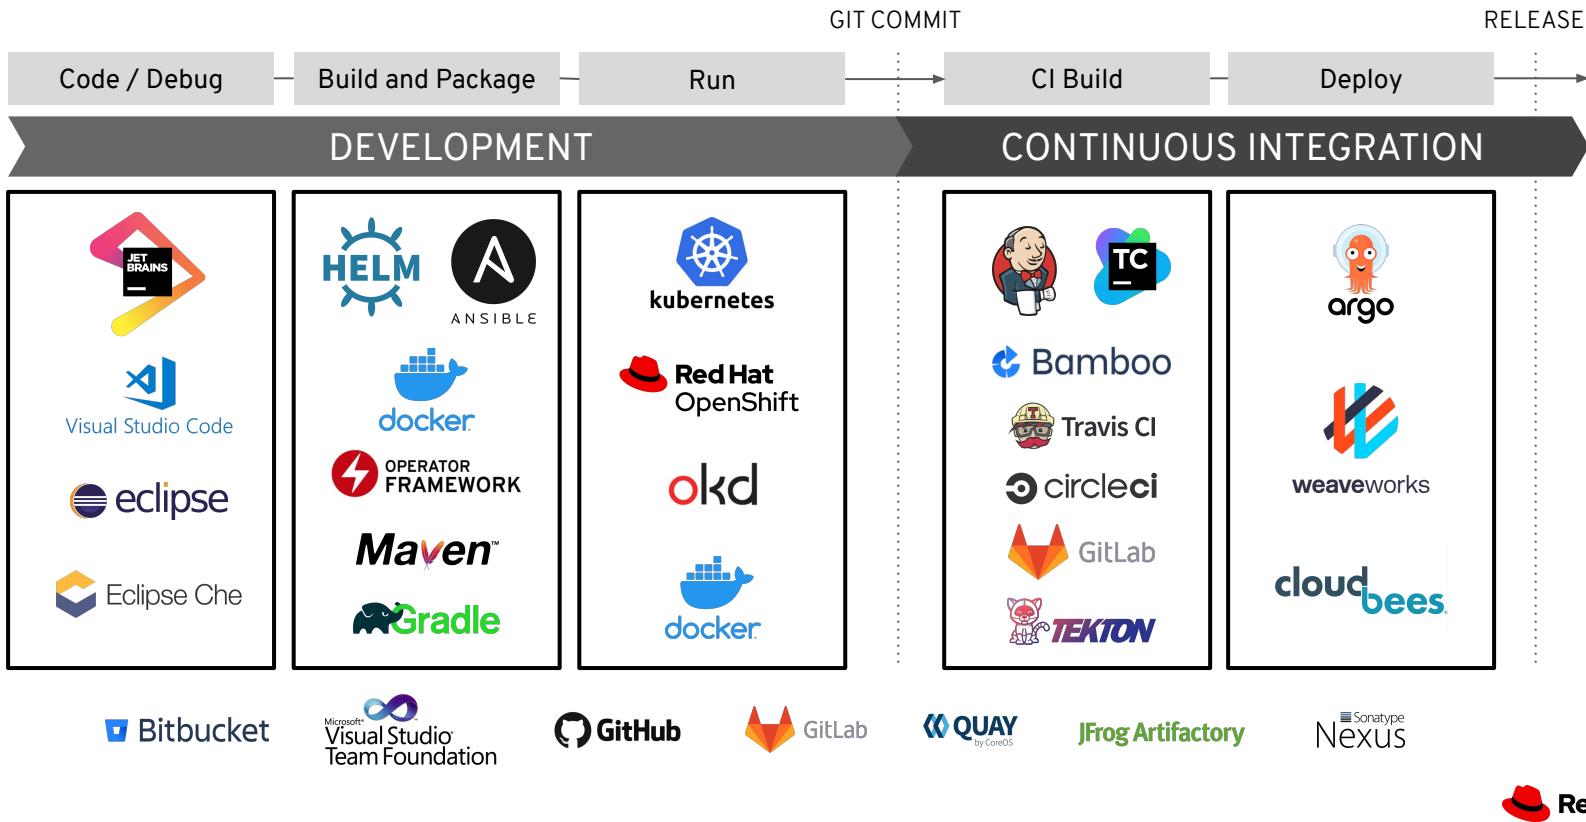
* **Develop on OpenShift for existing apps and new cloud-native apps**

* **Deploy on OpenShift in AWS, Azure, current environments or anywhere else**

One consistent development environment enables developers to move from project to project quickly. They don't need to learn new cloud-specific UIs and tools.

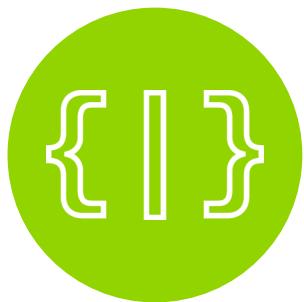
If teams need to leverage other cloud providers at any point it's easy to move specific apps because OpenShift provides an abstraction layer.

OpenShift integrates into your organization's preferred toolchain



BUILD & DEPLOY

BUILD AND DEPLOY CONTAINER IMAGES



DEPLOY YOUR
SOURCE CODE

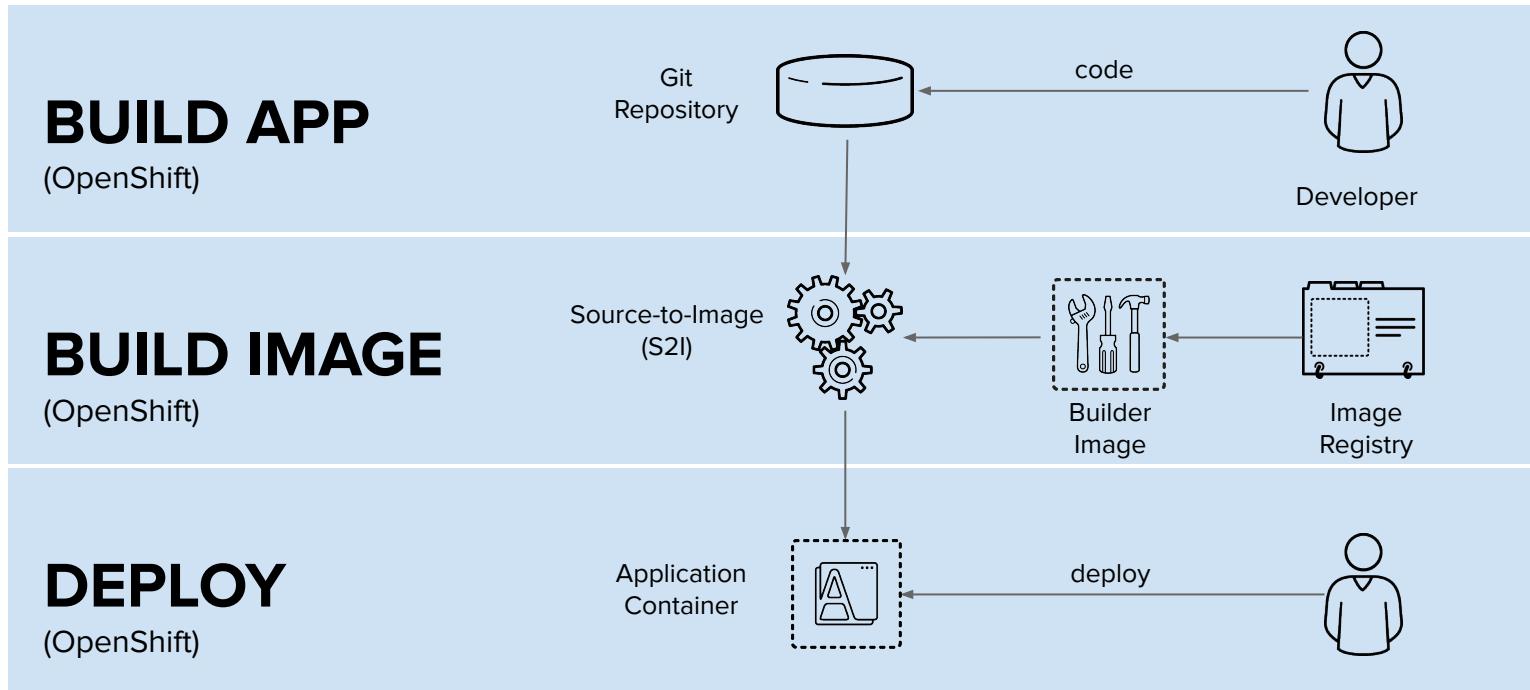


DEPLOY YOUR
APP BINARY

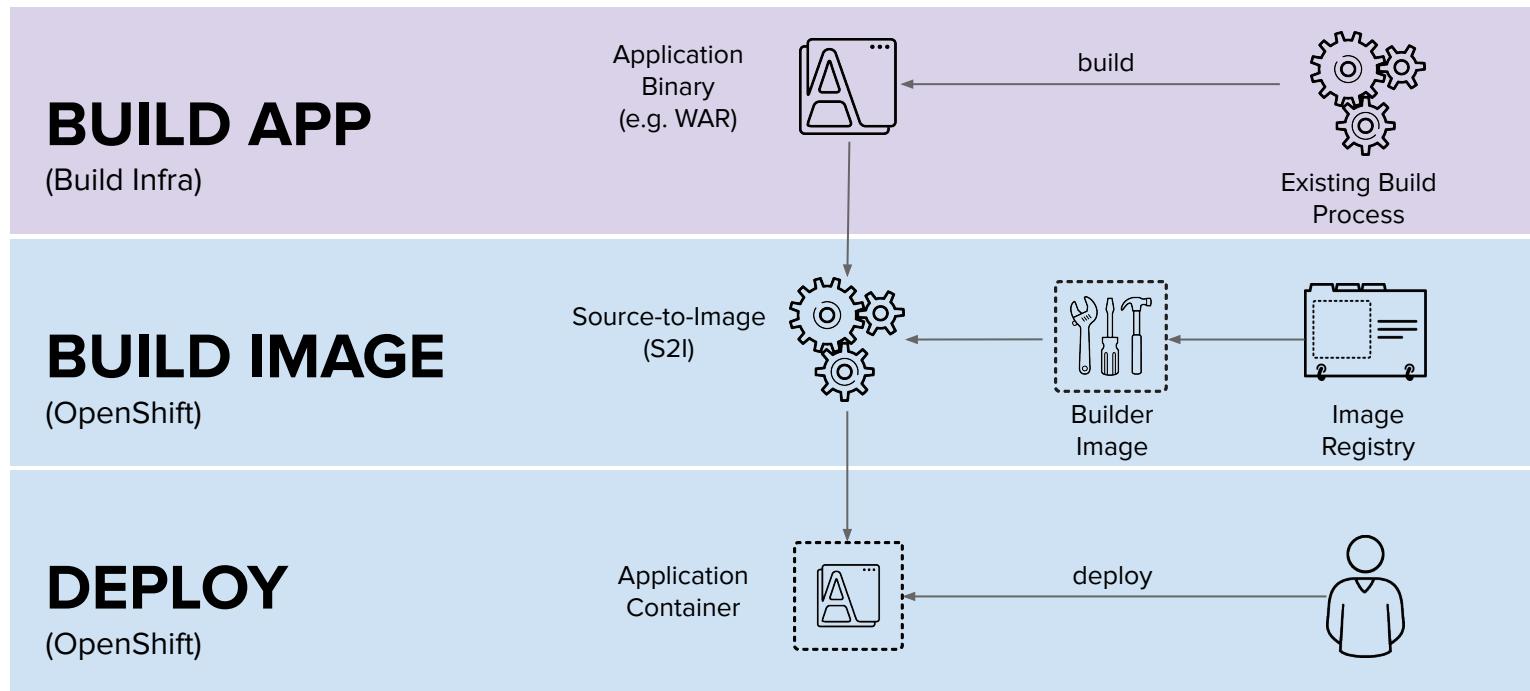


DEPLOY YOUR
CONTAINER IMAGE

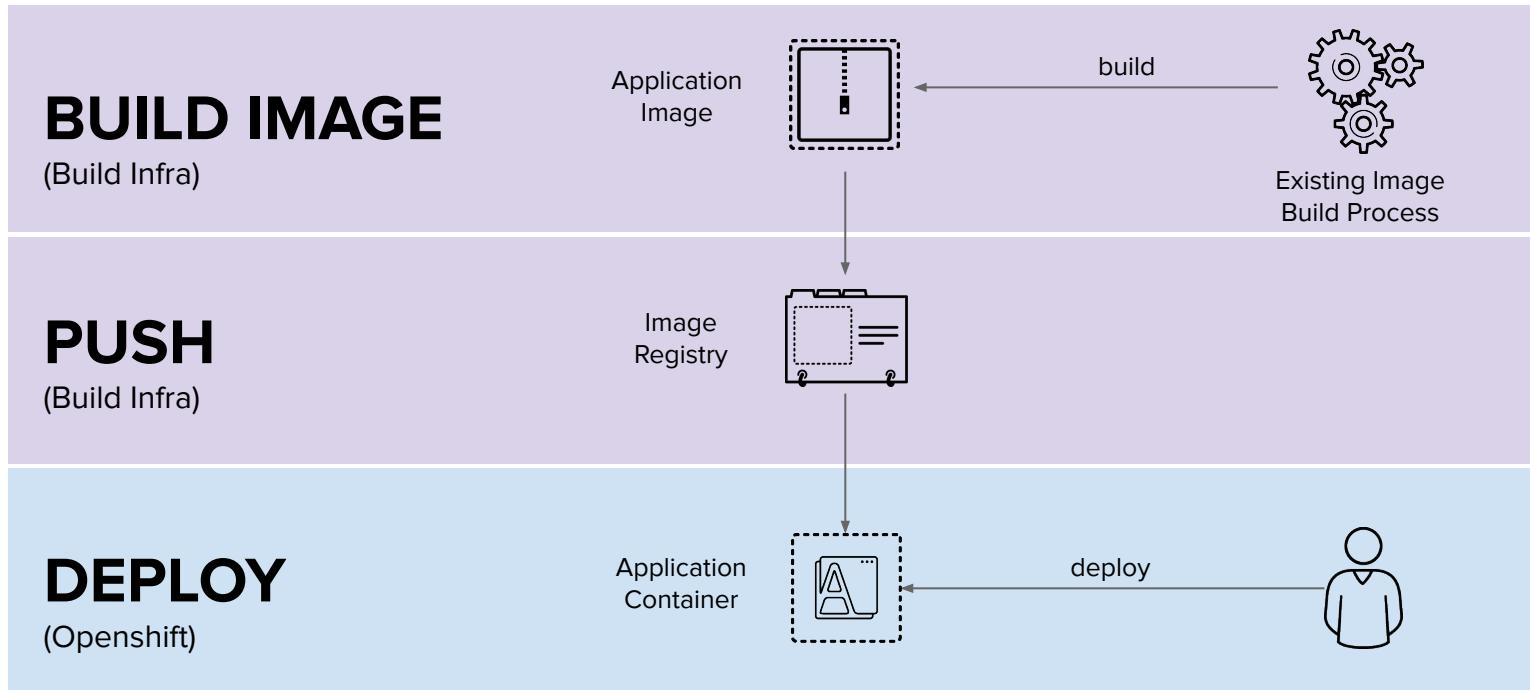
DEPLOY SOURCE CODE WITH SOURCE-TO-IMAGE (S2I)



DEPLOY APP BINARY WITH SOURCE-TO-IMAGE (S2I)

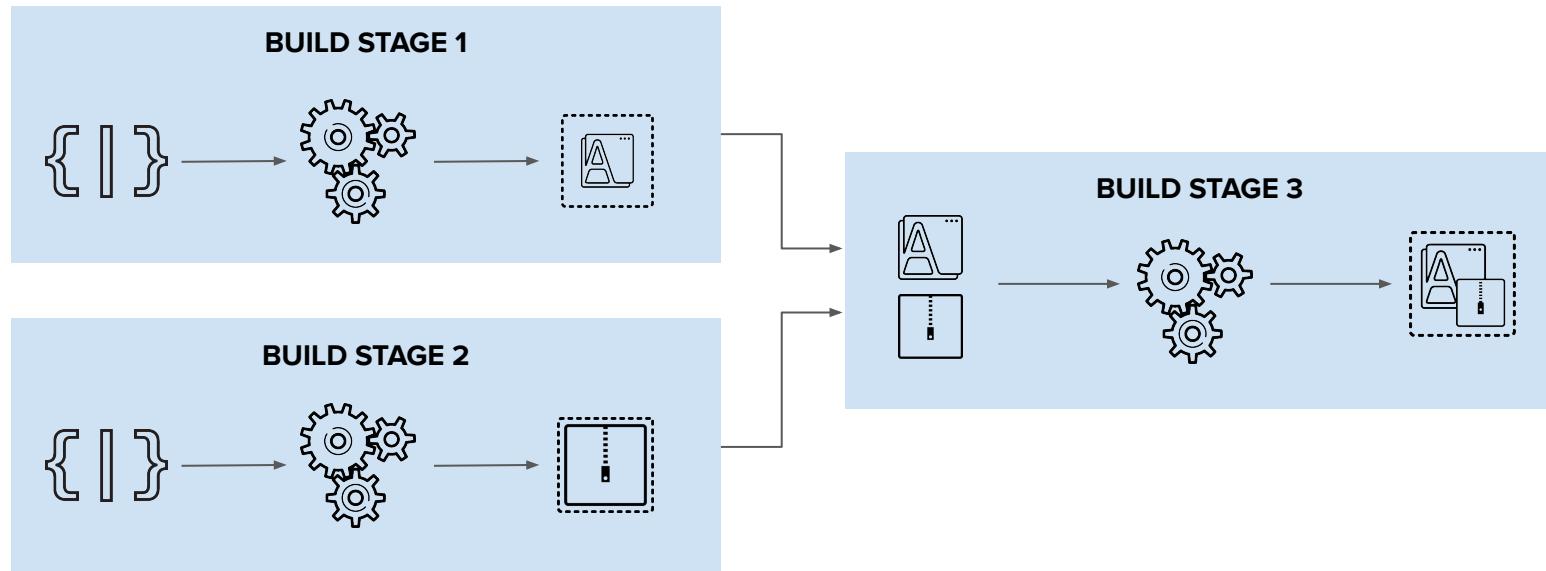


DEPLOY DOCKER IMAGE



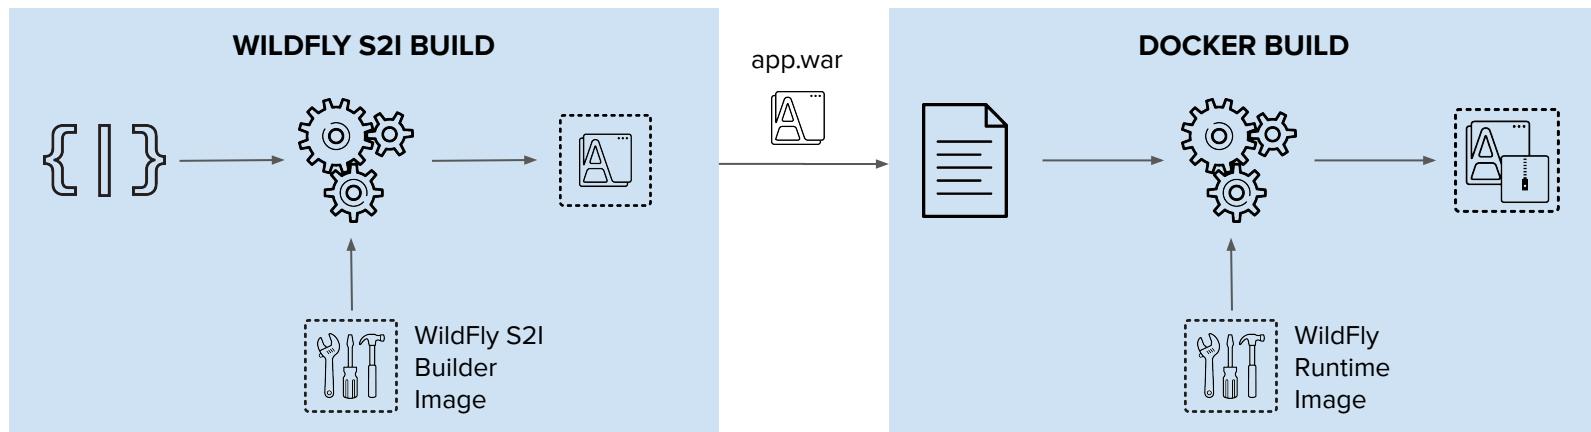
■ User/Tool Does ■ OpenShift Does

BUILD IMAGES IN MULTIPLE STAGES



EXAMPLE: USE ANY RUNTIME IMAGE WITH SOURCE-TO-IMAGE BUILDS

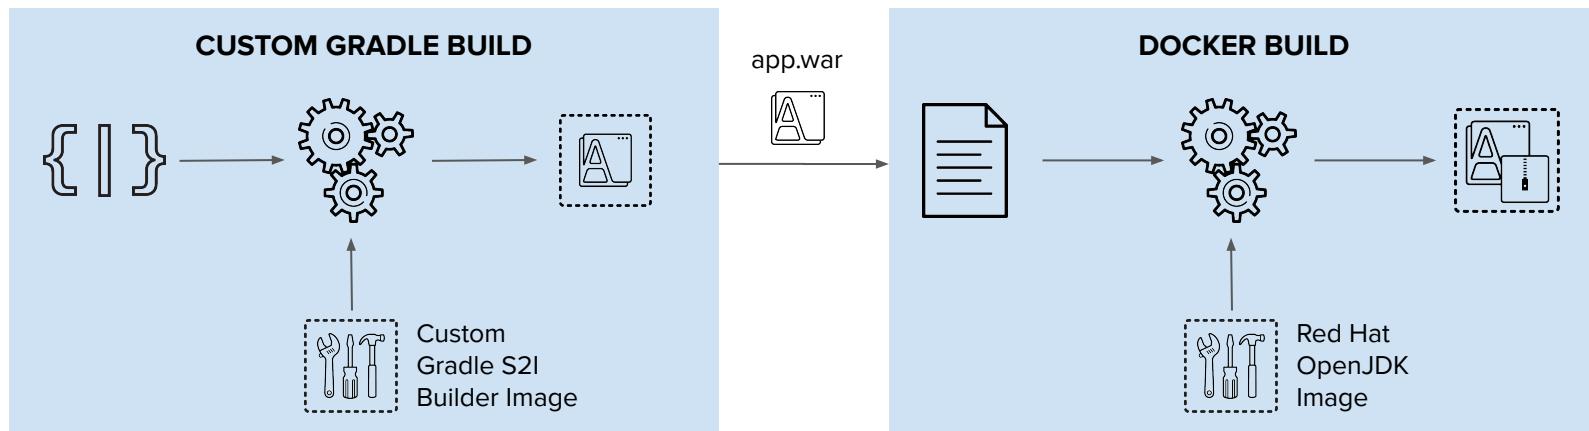
Use Source-to-Image to build app binaries and deploy on lean vanilla runtimes



read more on <https://blog.openshift.com/chaining-builds/>

EXAMPLE: USE ANY BUILD TOOL WITH OFFICIAL RUNTIME IMAGES

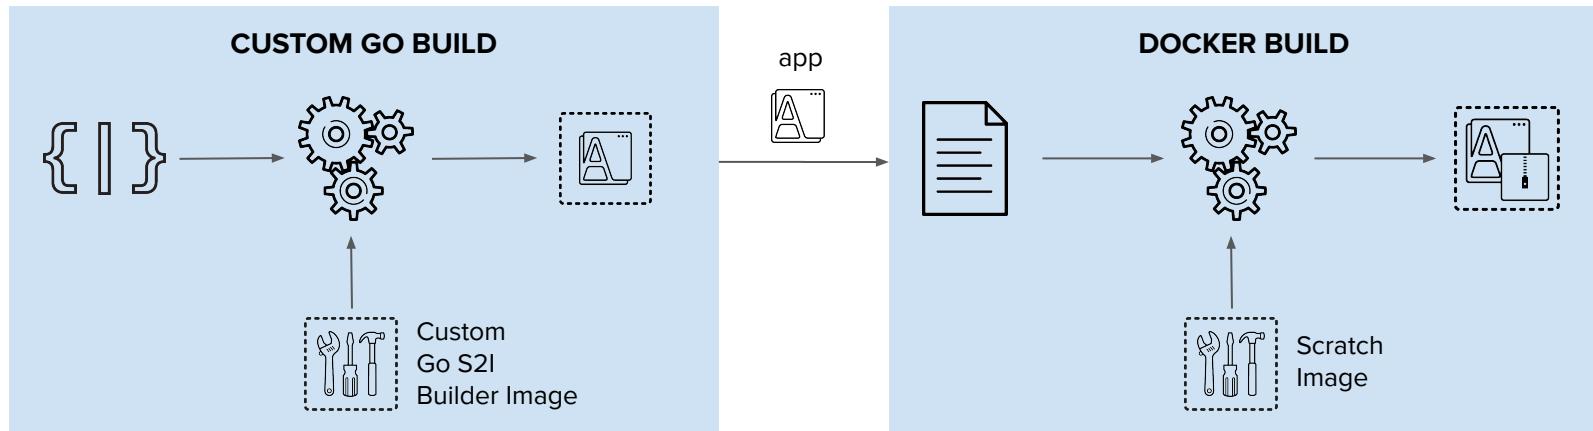
Use your choice of build tool like Gradle and deploy to official images like the JDK image



read more on <https://blog.openshift.com/chaining-builds/>

EXAMPLE: SMALL LEAN RUNTIMES

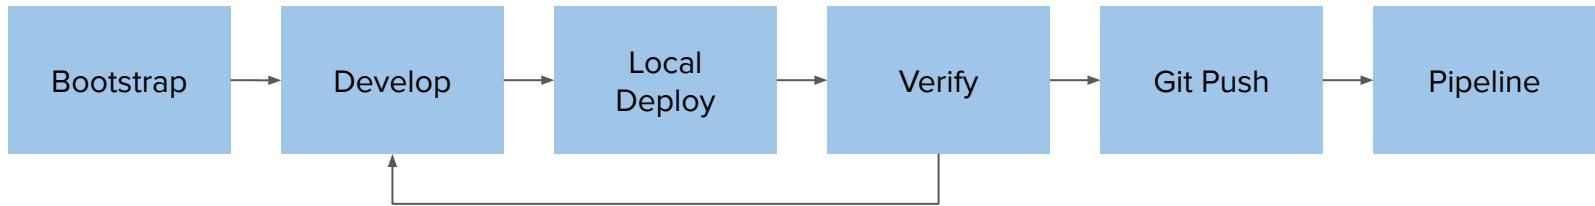
Build the app binary and deploy on small scratch images



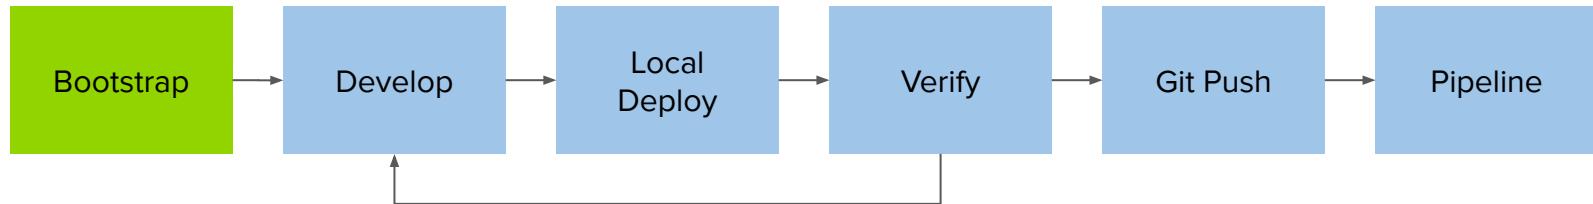
read more on <https://blog.openshift.com/chaining-builds/>

DEVELOPER WORKFLOW

LOCAL DEVELOPMENT WORKFLOW



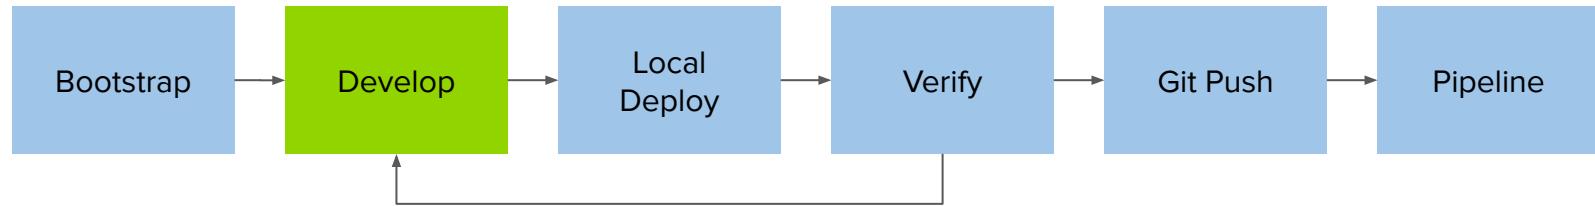
LOCAL DEVELOPMENT WORKFLOW



BOOTSTRAP

- Pick your programming language and application runtime of choice
- Create the project skeleton from scratch or use a generator such as
 - Maven archetypes
 - Quickstarts and Templates
 - OpenShift Generator
 - Spring Initializr

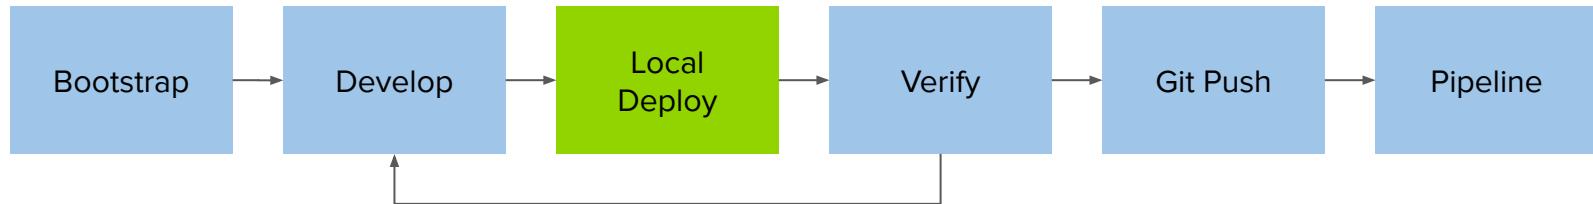
LOCAL DEVELOPMENT WORKFLOW



DEVELOP

- Pick your framework of choice such as Java EE, Spring, Ruby on Rails, Django, Express, ...
- Develop your application code using your editor or IDE of choice
- Build and test your application code locally using your build tools
- Create or generate OpenShift templates or Kubernetes objects

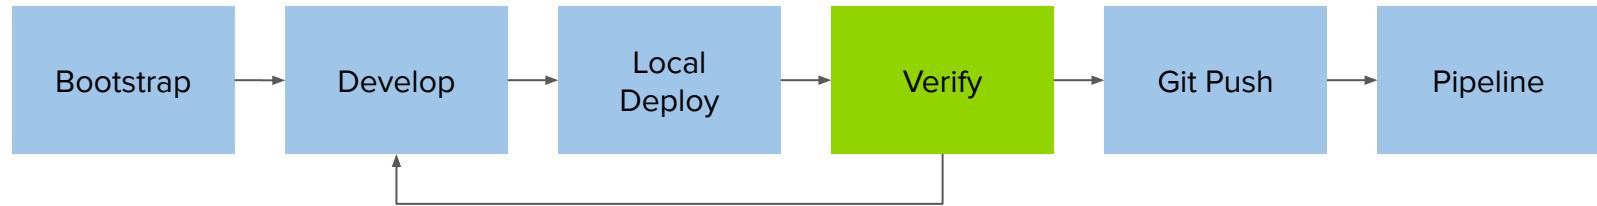
LOCAL DEVELOPMENT WORKFLOW



LOCAL DEPLOY

- Deploy your code on a local OpenShift cluster
 - Red Hat Container Development Kit (CDK), minishift and oc cluster
- Red Hat CDK provides a standard RHEL-based development environment
- Use binary deploy, maven or CLI rsync to push code or app binary directly into containers

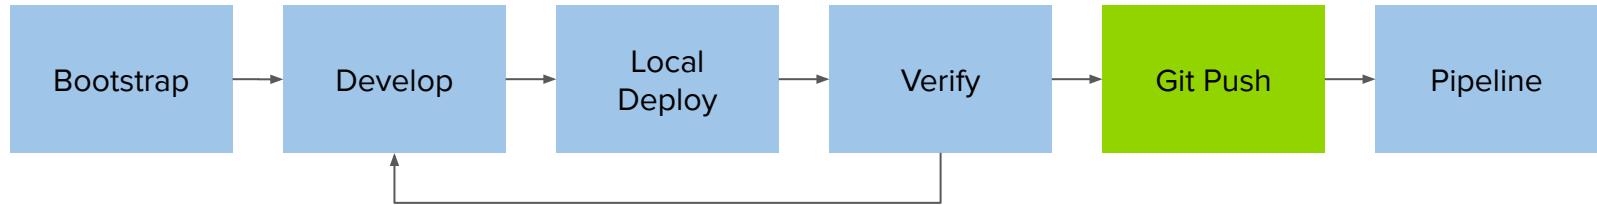
LOCAL DEVELOPMENT WORKFLOW



VERIFY

- Verify your code is working as expected
- Run any type of tests that are required with or without other components (database, etc)
- Based on the test results, change code, deploy, verify and repeat

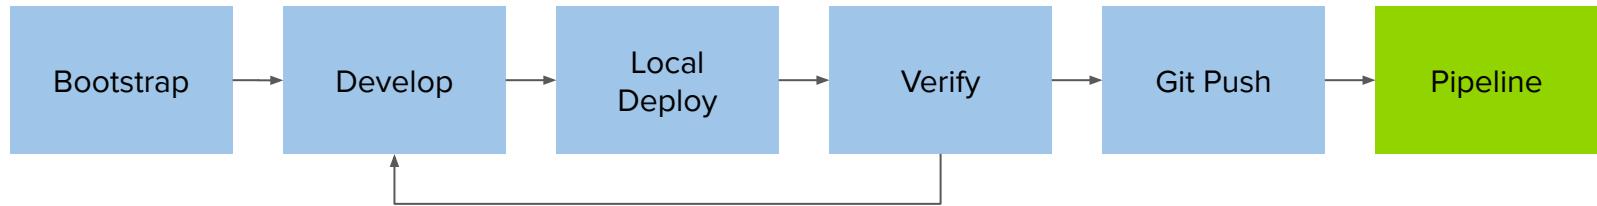
LOCAL DEVELOPMENT WORKFLOW



GIT PUSH

- Push the code and configuration to the Git repository
- If using Fork & Pull Request workflow, create a Pull Request
- If using code review workflow, participate in code review discussions

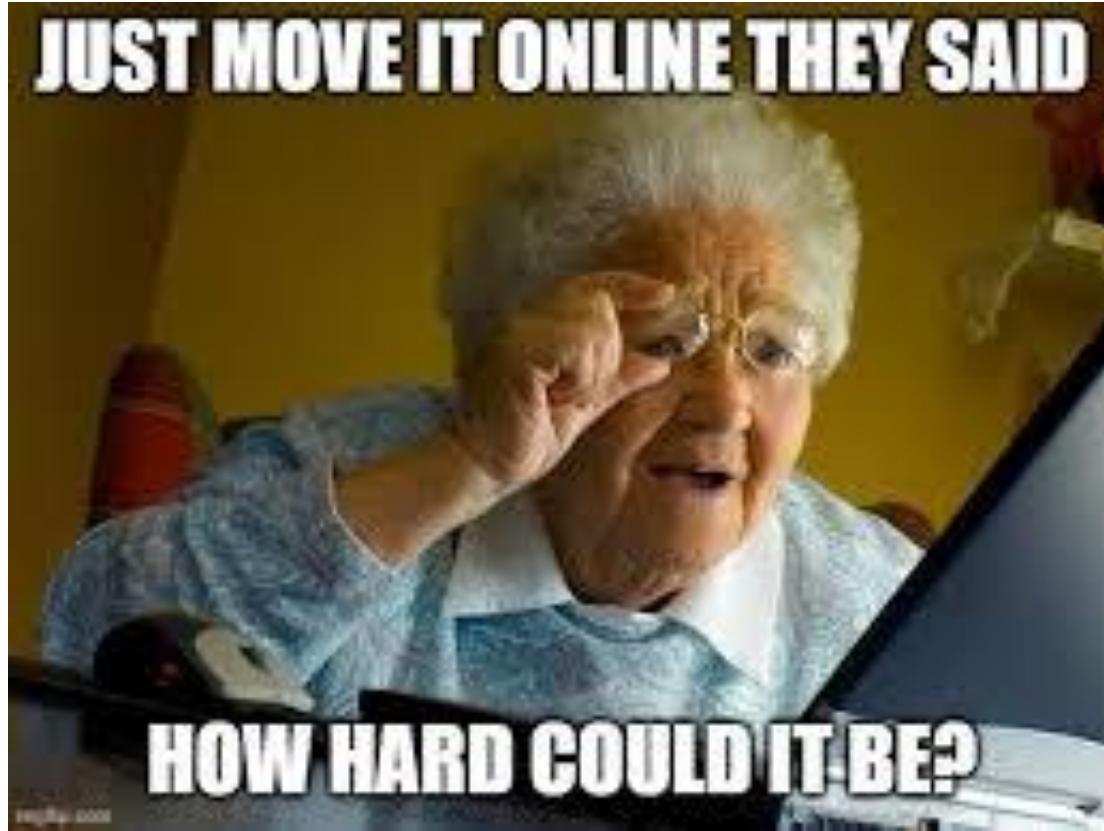
LOCAL DEVELOPMENT WORKFLOW



PIPELINE

- Pushing code to the Git repository triggers one or multiple deployment pipelines
- Design your pipelines based on your development workflow e.g. test the pull request
- Failure in the pipeline? Go back to the code and start again

Packaging Basics



What

- Now I have coded and my app works on my Kubernetes cluster
 - All fine
 - All done
- But wait...
- How to move those things from DEV to Test?
- How to release my software?
- (No, it's not just one image)
 - Deployment / DeploymentConfig
 - Service
 - PVCs
 - ConfigMaps
 - Route



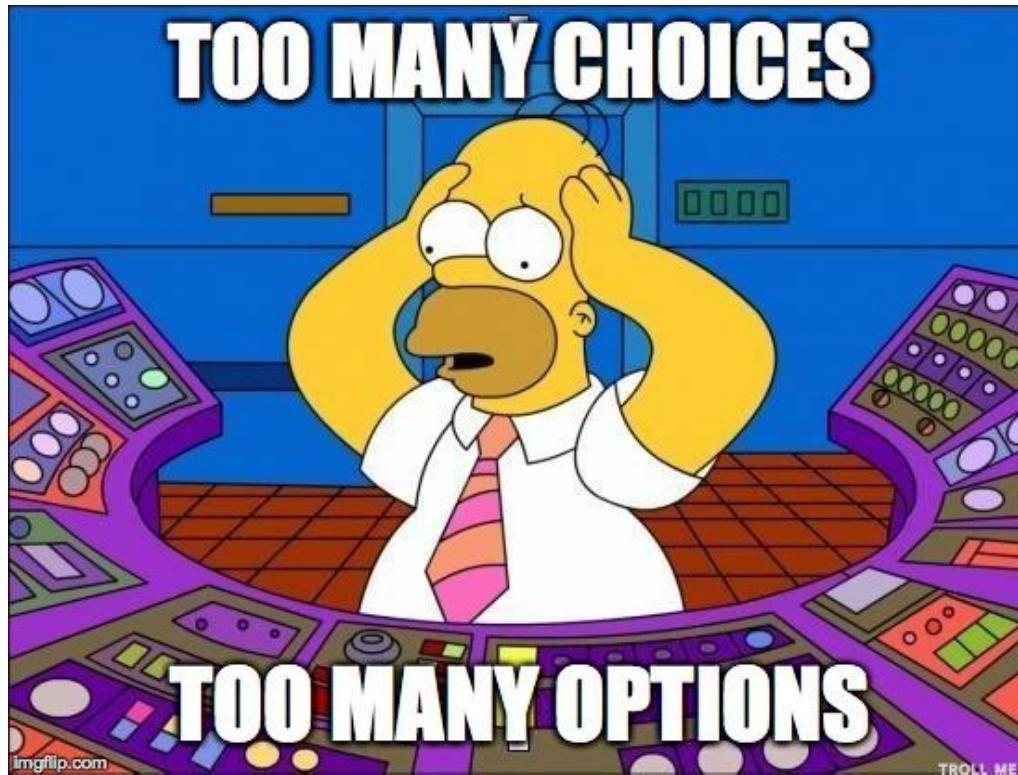
What

- How to automatically recreate your App with all resources and dependencies?
- Once you've created your App with all necessary resources, you need to somehow find a way to sync it with your stages (DEV/TEST/PRE-PROD...)
- How to redistribute your App?



Zip?
Tar?
Rsync?
Binaries?
Configuration?
Templates?
Helm Charts?
Operators?
DIY?
Kustomize?

Beer!



OpenShift Templates

What

- OpenShift Templates has been the first approach in OpenShift to make your project exportable and importable
 - All the RedHat middleware products were being used via templates
 - All Samples are still coming via templates
- Helm was too basic
- Helm required Tiller (a component which required root etc.)
- Operators did not exist
- It is basically a “`oc get is,bc,build,dc,is,rc,route,service -o yaml > test.yaml`”
 - With some editing
 - With QUITE some editing

Can I use it?

- Short answer: Of course!
- Longer answer: If you know that you are only using OpenShift inside your own Company then please use it.
 - A Template can configure even OpenShift specific types like BuildConfigs, ImageStreams, DeploymentConfigs...
 - Easily parametrized
 - With “oc process ...” there is an easy way to process a template
 - Can easily be used in your CI/CD pipelines
 - Can be versioned
 - Not too complex
 - Proven to work, easily understandable

Szenarios to use Templates

- In-Cluster movements (DEV → TEST)
- Cross-Cluster movements (TEST → PREPROD → PROD)
- Templating Application setups, including BC for use by others
- Building special examples as templates could be consumed via UI by developers
- Distributing your Apps

Drawbacks?

- Templates are OpenShift specific
- Templates are OpenShift specific
- Templates are OpenShift specific
- (And a template file could become VERY large as it contains ALL kubernetes yaml's required to run the app)

Template DEMO

kustomize.io

What

- Kustomize is a project originally founded Google
- It's in "kubectl apply -k" and "oc apply -k" now
- Has its own CLI interface, called kustomize
- It's NOT templating
- It's using overlays and patching

```
$ tree
.
├── base
│   ├── configMap.yaml
│   ├── deployment.yaml
│   ├── kustomization.yaml
│   ├── route.yaml
│   └── service.yaml
└── overlays
    ├── production
    │   ├── deployment.yaml
    │   └── kustomization.yaml
    └── staging
        ├── kustomization.yaml
        ├── map.yaml
        └── route.yaml
```

How it works

kustomization.yaml contains information about what to do and how

```
$ cat base/kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
metadata:
  name: arbitrary

# Example configuration for the webserver
# at https://github.com/monopole/hello
commonLabels:
  app: my-hello
  org: acmeCorporation

resources:
- deployment.yaml
- service.yaml
- configMap.yaml
- route.yaml
```

```
$ cat overlays/staging/kustomization.yaml
namePrefix: staging-
commonLabels:
  variant: staging
commonAnnotations:
  note: Hello, I am staging!
bases:
- ../../base
patchesStrategicMerge:
- map.yaml
- route.yaml
```

```
$ cat overlays/production/kustomization.yaml
namePrefix: production-
commonLabels:
  variant: production
commonAnnotations:
  note: Hello, I am production!
bases:
- ../../base
patchesStrategicMerge:
- deployment.yaml
- route.yaml
```

How it works

kustomize build or oc/kubectl apply -k does handle everything for you

```
$ kustomize build base
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: my-hello
    org: acmeCorporation
  name: the-service
spec:
  ports:
    - port: 8666
      protocol: TCP
      targetPort: 8080
  selector:
    app: my-hello
    deployment: hello
    org: acmeCorporation
  type: LoadBalancer
---
```

```
$ kustomize build overlays/staging
---
apiVersion: v1
kind: Service
metadata:
  annotations:
    note: Hello, I am staging!
  labels:
    app: my-hello
    org: acmeCorporation
    variant: staging
  name: staging-the-service
spec:
  ports:
    - port: 8666
      protocol: TCP
      targetPort: 8080
  selector:
    app: my-hello
    deployment: hello
    org: acmeCorporation
    variant: staging
  type: LoadBalancer
---
```

Can I use it?

- Short answer: Of course!
- Longer answer: If you're looking for a solution that integrates nicely with GitOps, you should definitely have a look
 - Kustomize can also configure OpenShift specific types like Routes etc.
 - It does NOT use parameters
 - With kustomize CLI there is a nice way to test your layers
 - Can easily be used in your CI/CD pipelines
 - Can be versioned
 - Not too complex
 - Proven to work, easily understandable

Scenarios to use kustomize

- In-Cluster movements (DEV → TEST)
- Cross-Cluster movements (TEST → PREPROD → PROD)
- GitOps
- OpenShift Pipelines / Tekton
- NOT usable for application publishing / distribution

Drawbacks?

- You can only change existing entries and add new ones...
- You can't use it for redistribution

Resources

[Automated Application Packaging and Distribution with OpenShift - Part 1/2 – Open Sourcerers](#)

<https://kustomize.io>

<https://github.com/kubernetes-sigs/kustomize>

[https://speakerdeck.com/spesnova/introduction-to-kus
tomize](https://speakerdeck.com/spesnova/introduction-to-kustomize)

<https://github.com/wpernath/kustomize-demo>

Kustomize-DEMO