

OpenShift Developer

Application Packaging

CI/CD

 [linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)

 [facebook.com/redhatinc](https://www.facebook.com/redhatinc)

 [youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)

 twitter.com/RedHat

Self introduction

Name: Wanja Pernath

Email: wpernath@redhat.com

Base: Germany (very close to the Alps)

Role: EMEA Technical Partner Development

Manager - OpenShift and MW

Experience: Years of Consulting, Training,

PreSales at Red Hat and before



Agenda / etc.

Agenda

- Application Packaging with OpenShift
 - Basics
 - OpenShift Templates
 - kustomize.io
 - Helm Charts
 - Operators
 - Summary
- Demo

Helm Charts

What

- Helm originally invented 2015 and introduced later that year at KubeCon
- Helm moved as Kubernetes subproject in 2016 as Helm 2.0
- Helm 3.x is now (since 2020) an official CNCF project
- Helm is THE package manager for Kubernetes Applications
 - Helm is like RPM / APT for Linux
 - Or maven / npm for Java / node.js
- Helm Charts can easily be created, installed into a Kubernetes Cluster and also being upgraded
- With the [Artifact Hub](#) you have a huge repository of available community driven and maintained charts for every need

Can I use it?

- Short answer: Of course, but mainly for distributing your app!
- Longer answer: If you have an app release and you have to make it available for others, create a Helm Chart for it and make it easy for your customers (regardless of internal or external) to consume it
 - If you are just looking for a way to move your app from one stage to the other, have a look at Templates or kustomize.io
 - Helm and ArtifactHub are a great resource to look at for components you might need

Szenarios to use Helm Charts

- Well used for distribution of Applications
- Internal distribution & external
- Not so great for use within CI/CD (but possible, of course)

Drawbacks?

- Learning curve of Helm Charts is steep at the beginning
- It adds another complexity to your app development cycle
- Client is a templating engine with its own DSL and complexity
- Helm is intended for Day-1 Operations
- Helm is intended for stateless application distribution

Helm 2 vs 3

- Helm 2 required a server component called Tiller
 - Another app on top of kubernetes which had to be managed and maintained
 - Tiller had its own RBAC and its own audit trail
 - Tiller was storing sensitive data in ConfigMaps
 - → Loss of visibility
- Helm 3 does not need a server side component
 - It uses native kubernetes approach and only a client side tool
- **⇒ This is the reason why OpenShift did not natively support Helm prior V3**

Resources

- [Helm.sh](#)
- [Spotlight on Helm](#)
- [To Helm or not?. Helm is becoming a very popular tool to... | by Stepan | FAUN](#)
- [From Templates to Openshift Helm Charts](#)
- [Working with Helm charts using the Developer perspective - Application life cycle management | Applications | OpenShift Container Platform 4.6](#)
- [How to make a Helm chart in 10 minutes](#)
- [Artifact Hub](#)
- <https://github.com/wpernath/helm-demo.git>
- [Automated Application Packaging And Distribution with OpenShift - Part 2/3 – Open Sourcerers](#)

Helm DEMO

Operators

What

- Operators were originally invented 2018 by coreos as a way to extend kubernetes by adding new custom resources and controllers or to help human operators to operate and manage stateful applications on kubernetes
- Operators are part of the kubernetes project
- Operators are like Helm Charts a way to distribute your app
- Operators contain all the domain logic required to manage the app
 - It understands how to scale up / down a stateful app
 - It understands how to do backups
- Operators are packaged as a mix of yaml definitions and a standard language like Go, Java etc.
- OperatorHub.io contains a nice set of available operators
- OperatorSDK helps you to create an own operator

Wait... stateless / stateful?

- Stateless
 - Kubernetes can manage stateless apps easily
 - Scaling is just a matter of adding a new pod
- Stateful
 - Imagine a mysql database
 - Scaling that up, means kubernetes is creating copies of the data
 - In fact, you then would have 3 different DBs
 - You need to find a way to properly scale mysql
 - Every app handles that differently (mysql vs. postgres vs. redis)
 - That domain logic needs to be put into an Operator to automate those tasks

Can I use it?

- Short answer: Of course!
- Longer answer:
 - Whenever you need to find a way to tell kubernetes how to manage your stateful complex app, you have to use Operators.
 - If your app is stateless or does not need special treatment, don't use Operators, think about Helm Charts then

Drawbacks?

- Quite complex
- You need to understand kubernetes properly

Resources

- [OperatorHub.io | The registry for Kubernetes Operators](#)
- [Operators on Red Hat OpenShift](#)
- [Operator SDK](#)
- [An intro to Kubernetes operators](#)
- [Kubernetes Operator simply explained in 10 mins](#)

Summary

Summary

- All 4 packaging mechanisms discussed are solving mainly 2 different use cases
 - Application Distribution
 - CI/CD
- Helm Charts, Kubernetes Operators and kustomize are standardized kubernetes or CNCF projects.
- Templates are OpenShift specific
- Unfortunately, you have to think about 2 different mechanisms in a typical project
 - You need CI/CD → kustomize or Templates
 - You might need app distribution → Helm or Operator

Summary - CI/CD

- Use OpenShift Templates if
 - You're purely on OpenShift
 - You need a quick and easy way to move your apps to other stages
 - You want to create special sample apps for developers
 - You want to be included in the developer perspective to choose from
 - You don't like the approach of kustomize (patch&merge)
- Use kustomize if
 - You just want to have a standard way of doing CI/CD
 - You don't like the template approach
 - You don't know if you're staying on OpenShift
 - You want to rely on kubernetes standards

Summary - Application Distribution

- Use Helm if
 - Your app is relatively easy and straight forward
 - Your app does not require special kubernetes configs
 - Your app is mainly a stateless application
- Use Operators if
 - Your app requires special handling, special kubernetes custom resources (CRDs)
 - Is complex and requires a special backup strategy
 - Needs several Dependencies
 - Have a special need for Day 2 Operations
 - Is a stateful application
- Good: You can even create Operators out of a Helm Chart

Resources

- [Kubernetes Operators and Helm — It takes Two to Tango](#)
- [Kubernetes Operators vs. Helm Charts: Which to Use and When](#)
- [Build Kubernetes Operators from Helm Charts in 5 steps](#)
- [Automated Application Packaging and Distribution with OpenShift – Part 1/2](#)
- [Automated Application Packaging And Distribution with OpenShift - Part 2/3](#)

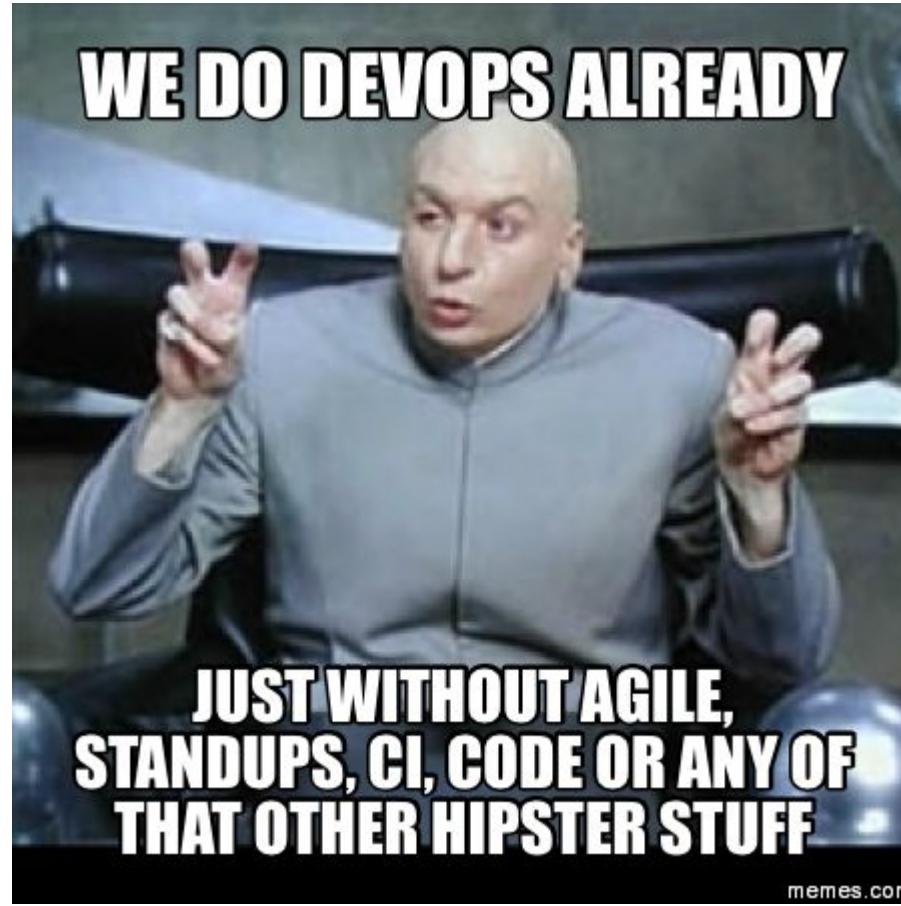
CI/CD with OpenShift

Basics

Wouldn't it be great, if everything could be automated?

Wouldn't it be great, if we simply do <insert hype here> and it solves all problems we have?

With DevOps we have a principle which solves most of the issues - but wait! The *DevOps Person* of the customer is currently on vacation



Basics - Developer

- Important: Separation of code and config!
- Writes the code of the App
- Writes set of build files (maven, gradle, etc.)
- Writes all needed tests (unit, integration, load)
- Writes Pipeline files
- Writes kubernetes manifest files
- Stores everything auditable in Git
- Finds a way and tools to combine all of the above

Basics - Operations

- Gets images, configs, test descriptions from Devs
- Adds own kubernetes manifests to the soup
- Makes sure, everything gets deployed
- Makes sure every dependency is installed and ready
- Thinks about security
- Thinks about networking
- Thinks about storage
- Thinks about compute power
- Thinks about plumbing everything together

Basics

CI/CD

=

Automation of Software Delivery

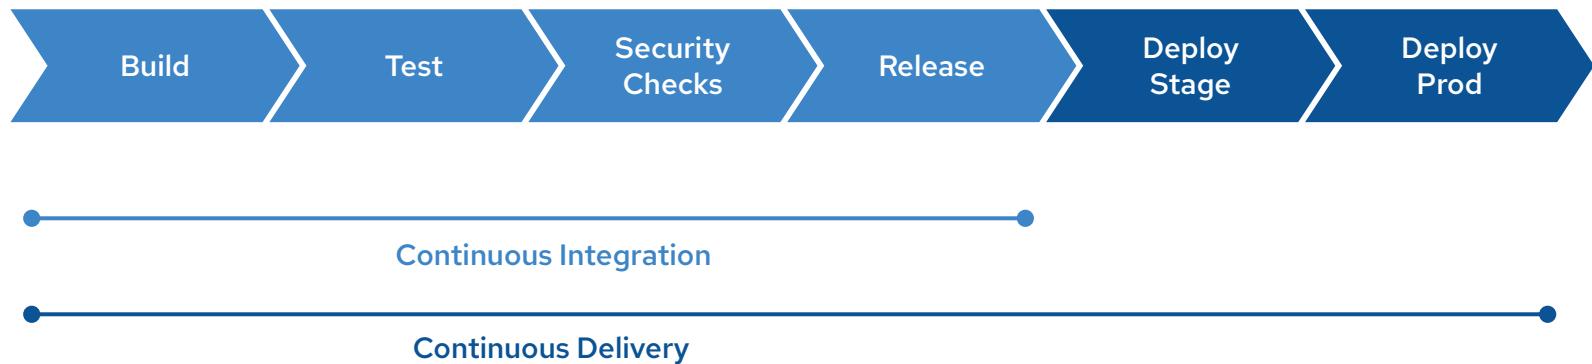
DevOps

=

Let's have Dev and Ops TALK to each other

Continuous Integration(CI) & Continuous Delivery (CD)

A key DevOps principle for automation, consistency and reliability



OpenShift Builds

Automate building
container images using
Kubernetes tools

OpenShift Builds



Kubernatives-native image build

A Kubernative-native way to building container images on OpenShift which is portable across Kubernetes distros



Supports multiple build strategies

Choose the build strategy that fits best your applications and skills: source-to-image, Dockerfile, and Cloud-Native Buildpacks

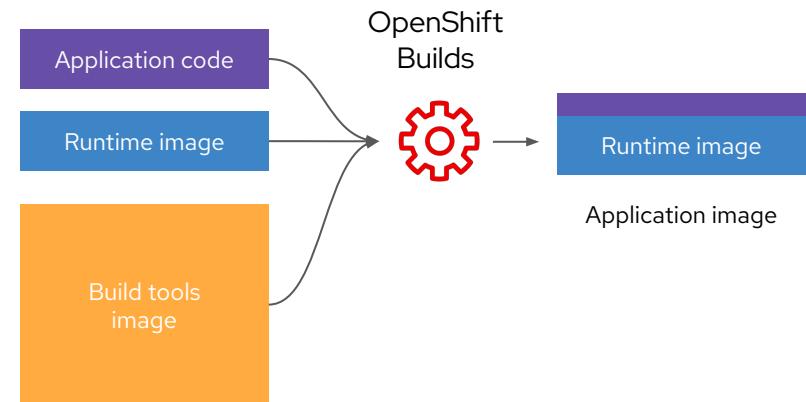


Extend with additional build strategies

Extend to use community Kubernetes builds strategies or your own custom builds

OpenShift Builds

- Build images on OpenShift and Kubernetes
- Use Kubernetes build tools
 - Source-to-Image
 - Buildpacks
 - Buildah
 - Kaniko
 - ...more
- Create lean application images
- Extend with your own build tools
- Based on Shipwright open-source project



OpenShift Builds

BuildStrategy

How to build images e.g. S2I, Buildpacks, etc

Build

What to build

BuildRun

Build execution details

OpenShift Builds

Cloud-Native Buildpacks	Source-to-Image (S2I)
<pre>kind: Build metadata: name: myapp-buildpack spec: source: url: https://github.com/myorg/myapp strategy: name: buildpacks-v3 builder: image: paketobuildpacks/builder:full output: image: quay.io/myorg/myapp:v1</pre>	<pre>kind: Build metadata: name: myapp-s2i spec: source: url: https://github.com/myorg/myapp strategy: name: source-to-image builder: image: registry.redhat.io/openjdk/openjdk-11-rhel8 output: image: quay.io/myorg/myapp:v1 runtime: image: docker.io/openjdk:11-jre-slim</pre>

Tekton / OpenShift Pipelines

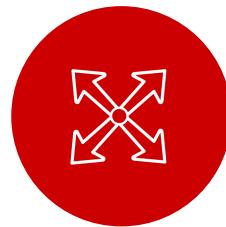
Kubernetes native on
demand delivery
pipelines

What is Cloud-Native CI/CD?



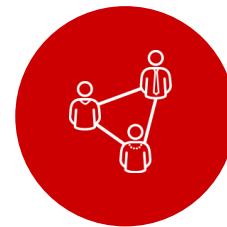
Containers

Built for container apps and runs on Kubernetes



Serverless

Runs serverless with no CI/CD engine to manage and maintain



DevOps

Designed with microservices and distributed teams in mind

Why Cloud-Native CI/CD?

Traditional CI/CD

Designed for Virtual Machines

Require IT Ops for CI engine maintenance

Plugins shared across CI engine

Plugin dependencies with undefined update cycles

No interoperability with Kubernetes resources

Admin manages persistence

Config baked into CI engine container

Cloud-Native CI/CD

Designed for Containers and Kubernetes

Pipeline as a service with no Ops overhead

Pipelines fully isolated from each other

Everything lifecycled as container images

Native Kubernetes resources

Platform manages persistence

Configured via Kubernetes ConfigMaps

Why Cloud-Native CI/CD?

Traditional CI/CD

Designed for Virtual Machines

Require IT Ops for CI engine maintenance



Plugins shared across CI engine
Jenkins
Plugins dependencies with undefined update cycles

No interoperability with Kubernetes resources

Admin manages persistence

Config baked into CI engine container

Cloud-Native CI/CD

Designed for Containers and Kubernetes

Pipeline as a service with no Ops overhead



Tasks fully isolated from each other
TEKTON
Everything literally has a Docker image

Native Kubernetes resources

Platform manages persistence

Configured via Kubernetes ConfigMaps

OpenShift Pipelines



Built for Kubernetes

Cloud-native pipelines taking advantage of Kubernetes execution and , operational model and concepts



Scale on-demand

Pipelines run and scale on-demand in isolated containers, with repeatable and predictable outcomes



Secure pipeline execution

Kubernetes RBAC and security model ensures security consistently across pipelines and workloads



Flexible and powerful

Granular control over pipeline execution details on Kubernetes, to support your exact requirements



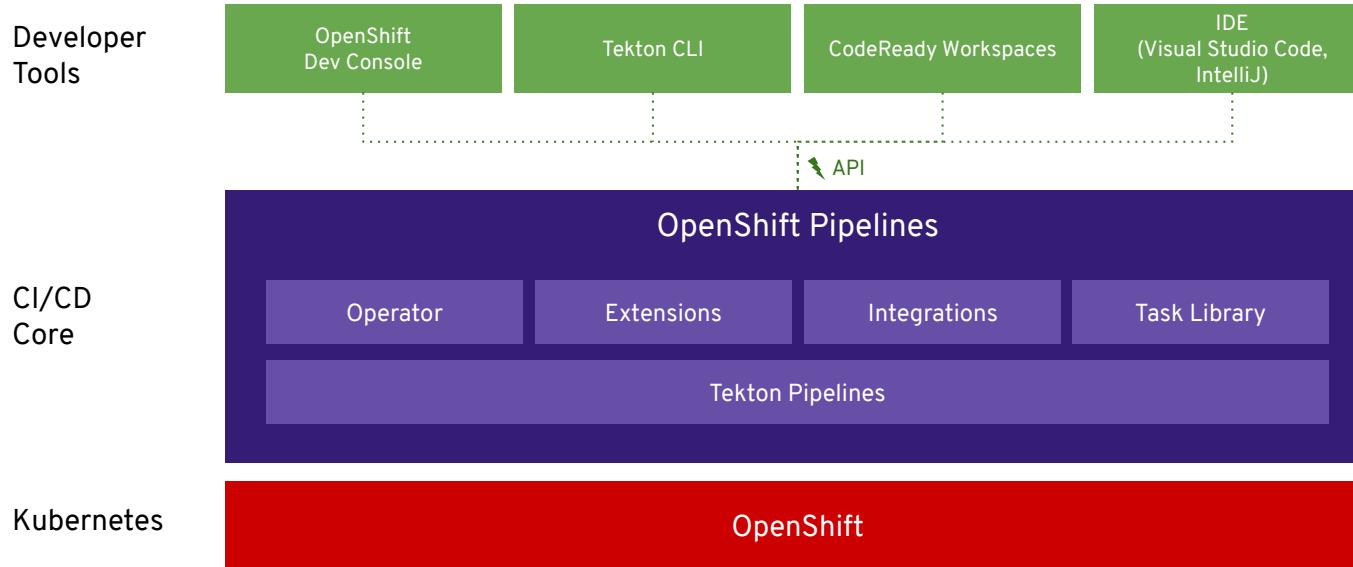
An open-source project for providing a set of shared and standard components for building Kubernetes-style CI/CD systems

41



Governed by the Continuous Delivery Foundation
Contributions from Google, Red Hat, Cloudbees, IBM, Pivotal and many more

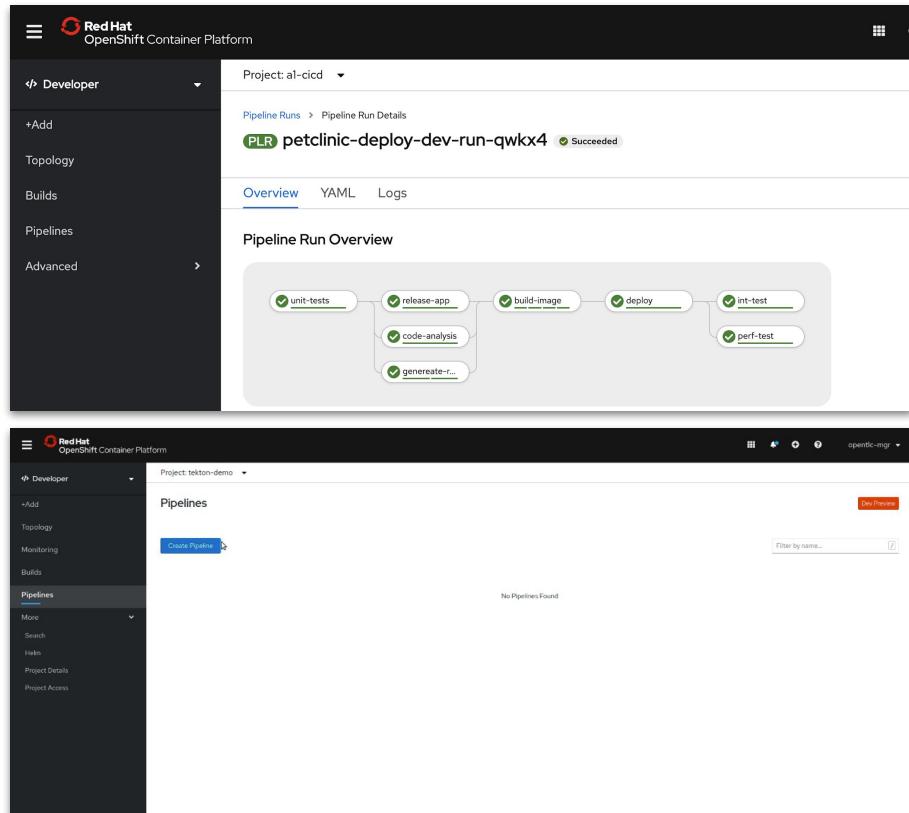
OpenShift Pipelines



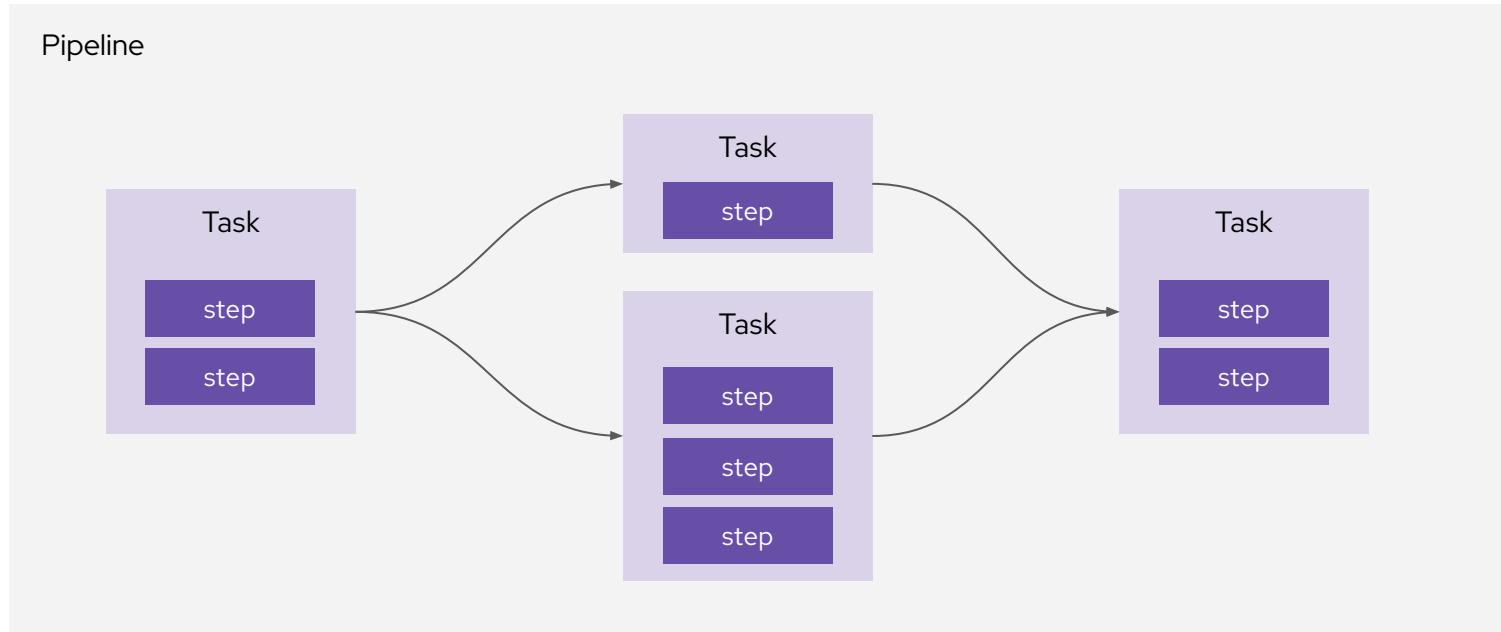
OpenShift Pipelines

- Based on Tekton Pipelines
- Kubernetes-native declarative CI/CD
- Pipelines run on-demand in isolated containers
- No central server to maintain! No plugin conflicts!
- Task library and integration with Tekton Hub
- Secure pipelines aligned with Kubernetes RBAC
- Visual and IDE-based pipeline authoring
- Pipeline templates when importing apps
- Automated install and upgrades via OperatorHub
- CLI, Web, VS Code and IntelliJ plugins

43



Tekton Concepts



Tekton Concepts: step

- Run command or script in a container
- Kubernetes container spec
 - Env vars
 - Volumes
 - Config maps
 - Secrets

```
- name: build
  image: maven:3.6.0-jdk-8-slim
  command: ["mvn"]
  args: ["install"]
```

```
- name: parse-yaml
  image: python3
  script: |-
    #!/usr/bin/env python3
    ...
```

Tekton Concepts: Task

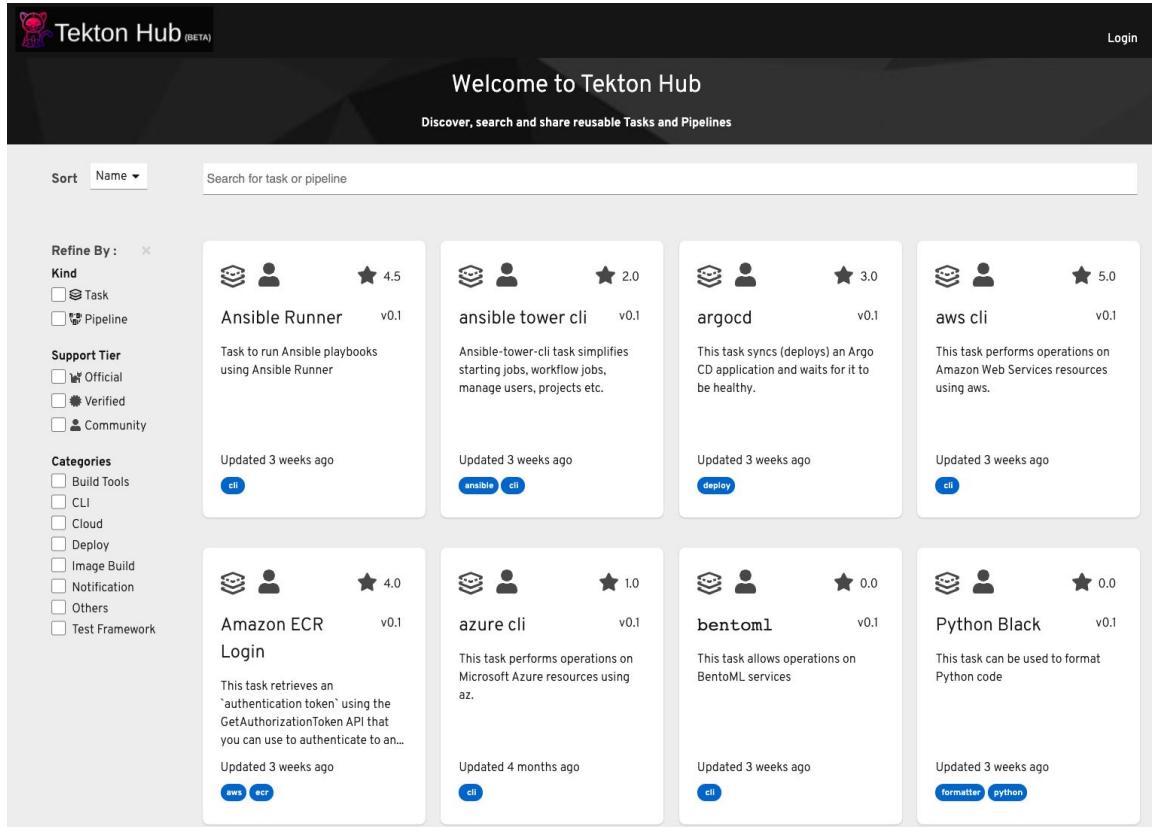
- Performs a specific task
- List of steps
- Steps run sequentially
- Reusable

```
kind: Task
metadata:
  name: buildah
spec:
  params:
    - name: IMAGE
  steps:
    - name: build
      image: quay.io/buildah/stable:latest
      command: ["buildah"]
      args: ["bud", ".", "-t", "$(params.IMAGE)"]
    - name: push
      image: quay.io/buildah/stable:latest
      script: |
        buildah push $(params.IMAGE) docker://$(params.IMAGE)
```

Tekton Hub

Search, discover and install Tekton Tasks

47



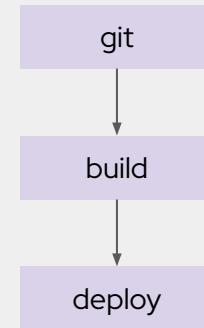
The screenshot shows the Tekton Hub interface, a platform for discovering and sharing reusable Tasks and Pipelines. The top navigation bar includes the Tekton Hub logo, a search bar, and a 'Login' button. The main content area features a search bar and a 'Sort' dropdown set to 'Name'. A sidebar on the left provides filtering options for 'Kind' (Task, Pipeline), 'Support Tier' (Official, Verified, Community), and 'Categories' (Build Tools, CLI, Cloud, Deploy, Image Build, Notification, Others, Test Framework). The main area displays a grid of 12 Task cards, each with a star rating, version, description, and update history. The tasks shown are Ansible Runner (v0.1, 4.5 stars), ansible tower cli (v0.1, 2.0 stars), argocd (v0.1, 3.0 stars), aws cli (v0.1, 5.0 stars), Amazon ECR Login (v0.1, 4.0 stars), azure cli (v0.1, 1.0 stars), bentoml (v0.1, 0.0 stars), and Python Black (v0.1, 0.0 stars).

Task	Version	Rating	Description	Updated	Tags
Ansible Runner	v0.1	4.5	Task to run Ansible playbooks using Ansible Runner	3 weeks ago	cli
ansible tower cli	v0.1	2.0	Ansible-tower-cli task simplifies starting jobs, workflow jobs, manage users, projects etc.	3 weeks ago	ansible, cli
argocd	v0.1	3.0	This task syncs (deploys) an Argo CD application and waits for it to be healthy.	3 weeks ago	deploy
aws cli	v0.1	5.0	This task performs operations on Amazon Web Services resources using aws.	3 weeks ago	cli
Amazon ECR Login	v0.1	4.0	This task retrieves an 'authentication token' using the GetAuthorizationToken API that you can use to authenticate to an...	3 weeks ago	aws, ecr
azure cli	v0.1	1.0	This task performs operations on Microsoft Azure resources using az.	4 months ago	cli
bentoml	v0.1	0.0	This task allows operations on BentoML services	3 weeks ago	cli
Python Black	v0.1	0.0	This task can be used to format Python code	3 weeks ago	formatter, python

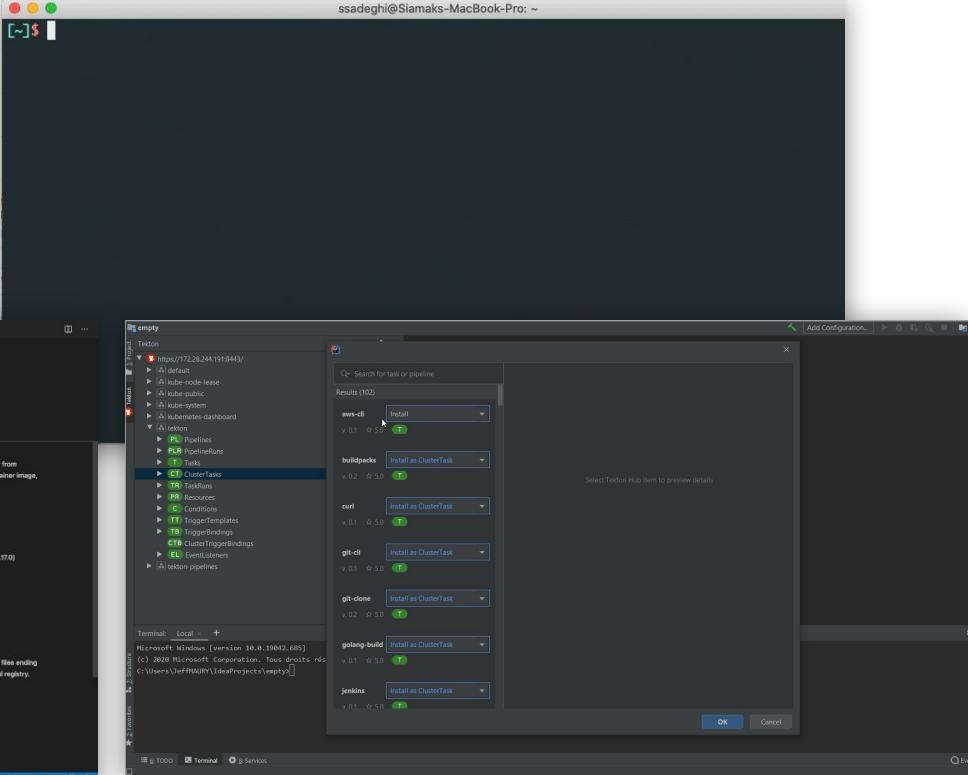
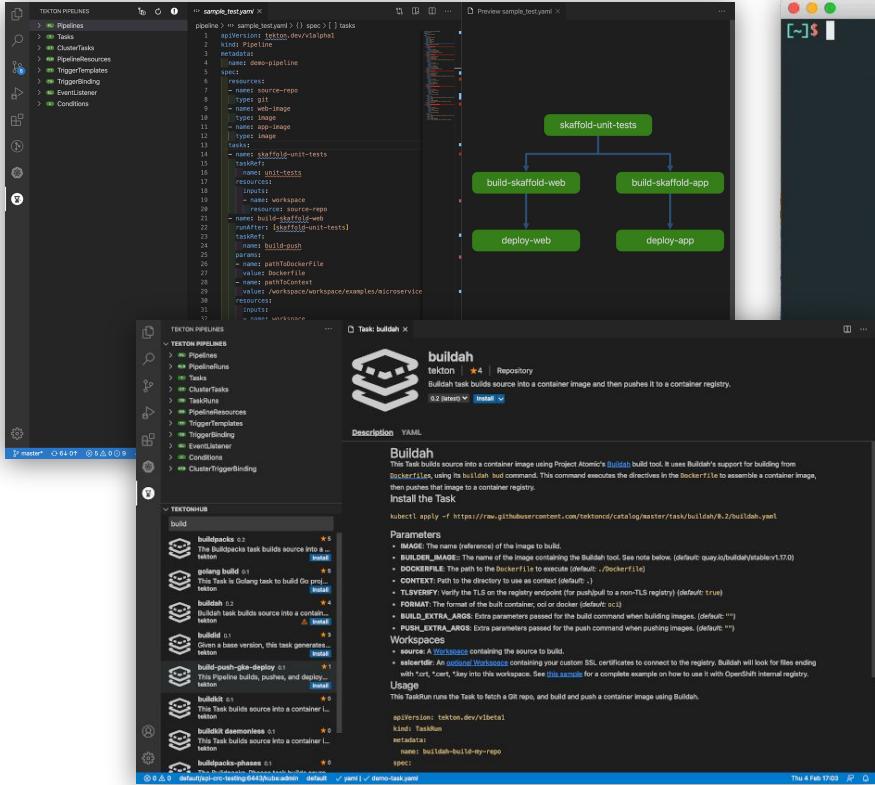
Tekton Concepts: Pipeline

- A graph of Tasks: concurrent & sequential
- Tasks run on different nodes
- Task execution logic
 - Conditional
 - Retries
- Share data between tasks

```
kind: Pipeline
metadata:
  name: deploy-dev
spec:
  params:
    - name: IMAGE_TAG
  tasks:
    - name: git
      taskRef:
        name: git-clone
      params: [...]
    - name: build
      taskRef:
        name: maven
      params: [...]
      runAfter: ["git"]
    - name: deploy
      taskRef:
        name: knative-deploy
      params: [...]
      runAfter: ["build"]
```



Tekton CLI, Visual Studio Code, and IntelliJ



DEMO TIME

<https://github.com/wpernath/openshift-cicd-demo>

OpenShift GitOps

Declarative GitOps for
multi-cluster continuous
delivery

What is GitOps?

GitOps is when the infrastructure and / or application state is fully represented by the contents of a git repository. Any changes to the repository are reflected in the corresponding state of the associated infrastructure and applications through automation

GitOps is a natural evolution of Agile and DevOps (and Kubernetes) methodologies

Why GitOps

It takes too long to provision a new environment!

The app behaves different in prod than in test!

I have no visibility or record of config changes in deployments!

I can't easily rollback changes to a specific version

Environments are all manually configured!!!

Production deployments have a low success rate!

I can't audit config changes!!



GitOps Benefits

- All changes are auditable
- Standard roll-forward or backwards in the event of failure
- Disaster recovery is “reapply the current state of the manifests”
- Experience is “pushes and pull-requests”

OpenShift and GitOps - A great match

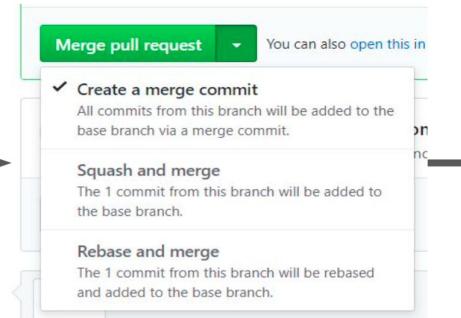
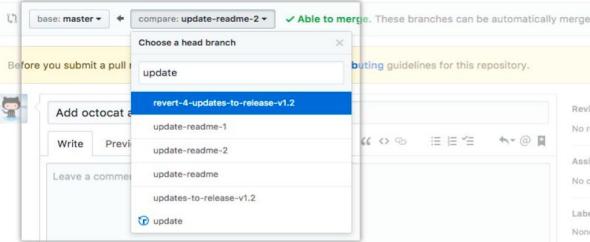
- OpenShift is a declarative environment
 - Cluster configuration is declared and Operators make it happen
 - Application deployments are declared and Kubernetes scheduler makes it happen
- GitOps in traditional environments requires automation/scripting, declarative environment minimizes or eliminates this need
- Declarations are yaml files which are easily stored and managed in git



Day 2 operations: All changes triggered from Git

Open a pull request

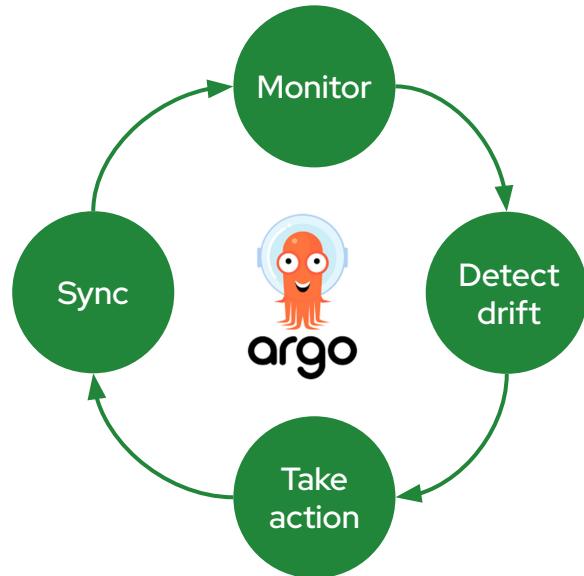
Create a new pull request by comparing changes across two branches. If you need to, you can also compare across forks.



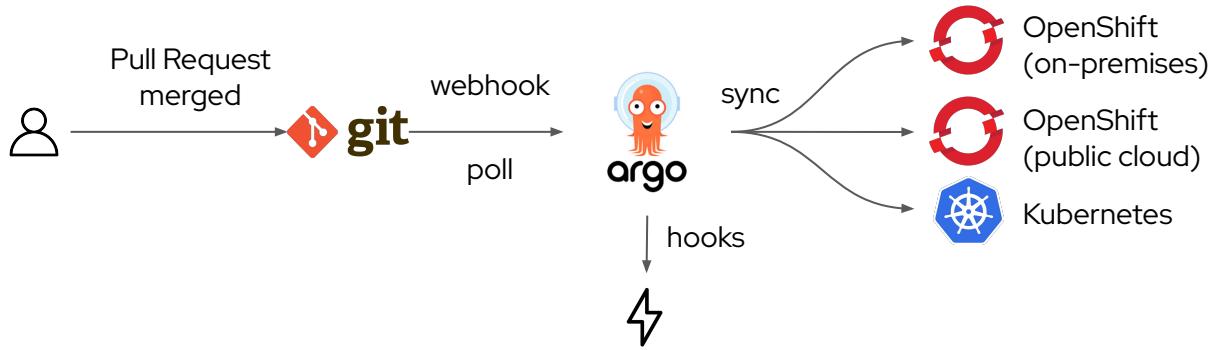
```
$ tkn pipelinerun logs update-from-master-run-g6s45
[run-kubectl] {"level": "info", "ts": 1580989837.3045664, "logger": "fallback-logger", "caller": "logging/config.go:69", "msg": "Fetch GitHub commit ID from kodata failed: \"\\$KO_DATA_PATH\" does not exist or is empty"}
[run-kubectl] serviceaccount/demo-sa unchanged
[run-kubectl] clusterrolebinding.rbac.authorization.k8s.io/tekton-triggers-openshift-binding unchanged
[run-kubectl] eventlistener.tekton.dev/demo-event-listener configured
[run-kubectl] task.tekton.dev/deploy-from-source-task configured
[run-kubectl] triggerbinding.tekton.dev/update-from-master-binding unchanged
[run-kubectl] triggertemplate.tekton.dev/update-from-master-template unchanged
```

Argo CD

- Cluster and application configuration versioned in Git
- Automatically syncs configuration from Git to clusters
- Drift detection, visualization and correction
- Granular control over sync order for complex rollouts
- Rollback and rollforward to any Git commit
- Manifest templating support (Helm, Kustomize, etc)
- Visual insight into sync status and history



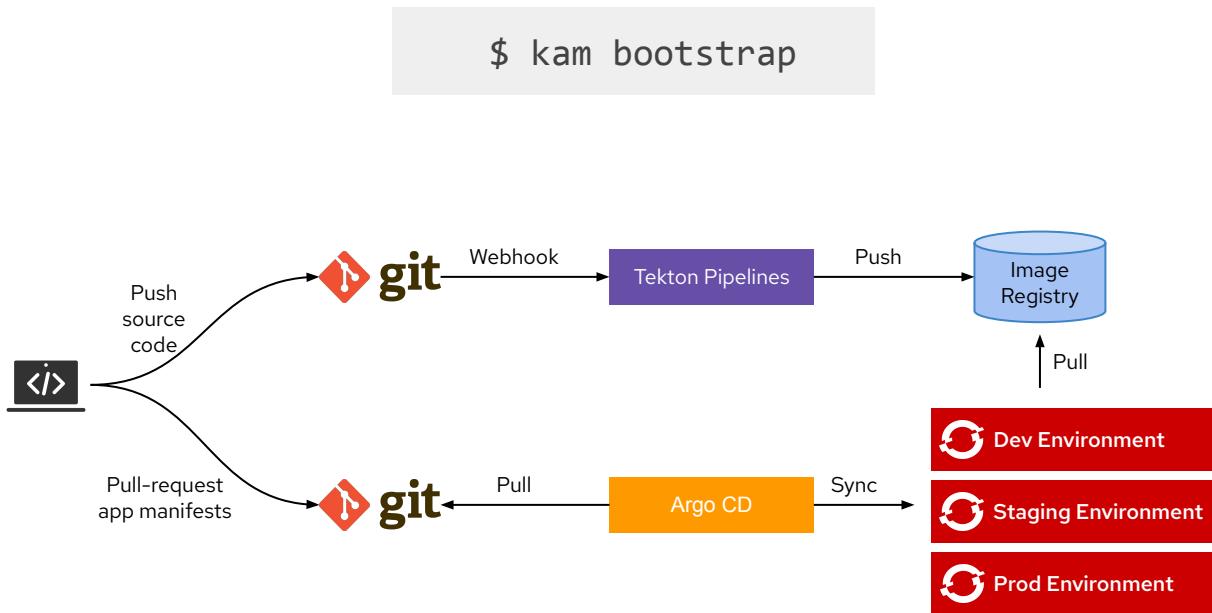
Argo CD



The screenshot shows the Argo CD web interface with the following details:

- Header:** Applications / realtime
- Toolbar:** APP DETAILS, APP DIFF, SYNC, SYNC STATUS, HISTORY AND ROLLBACK, DELETE, REFRESH
- Status:** Healthy
- Sync Status:** Synced to latest 691471, Authored by Geert Baekie <geertbaekie@MacBook...
- Synchronization:** Succeeded 38 minutes ago (Wed Dec 25 2019 13:08:45 GMT+0100) by Geert Baekie <geertbaekie@MacBook.local> using 1.0.5
- Deployment:** realtimeapp (rev.2) replicant realtimeapp-5bcb859455 (rev.2) replicant realtimeapp-6f449c7f6b (rev.1) replicant
- Redis Application:** redisapp (rev.1) deployment redisapp-74b945bf (rev.1) replicant redisapp-74b945bf-r6jj (running 1/1 pod)

GitOps Application Manager CLI



OpenShift GitOps



Multi-cluster config management

Declaratively manage cluster and application configurations across multi-cluster OpenShift and Kubernetes infrastructure with Argo CD



Automated Argo CD install and upgrade

Automated install, configurations and upgrade of Argo CD through OperatorHub



Opinionated GitOps bootstrapping

Bootstrap end-to-end GitOps workflows for application delivery using Argo CD and Tekton with GitOps Application Manager CLI



Deployments and environments insights

Visibility into application deployments across environments and the history of deployments in the OpenShift Console

ArgoCD - Challenges

- Repo structure for manifests
 - One Repo or
 - Separate Repos for environments
- Order dependent deployments
- Non-declarative deployments
- Integration with CI/CD tools (Jenkins, Pipelines...)
 - Who does manage deployments?



Multiple repositories

/taxi.git

deploy

pipelines

pkg/cmd/booktaxi

web

Dockerfile

/taxi-config-stage.git

deploy

pipelines

README.md

/taxi-config-prod.git

deploy

pipelines

README.md

/taxi-config-test.git

deploy

pipelines

README.md

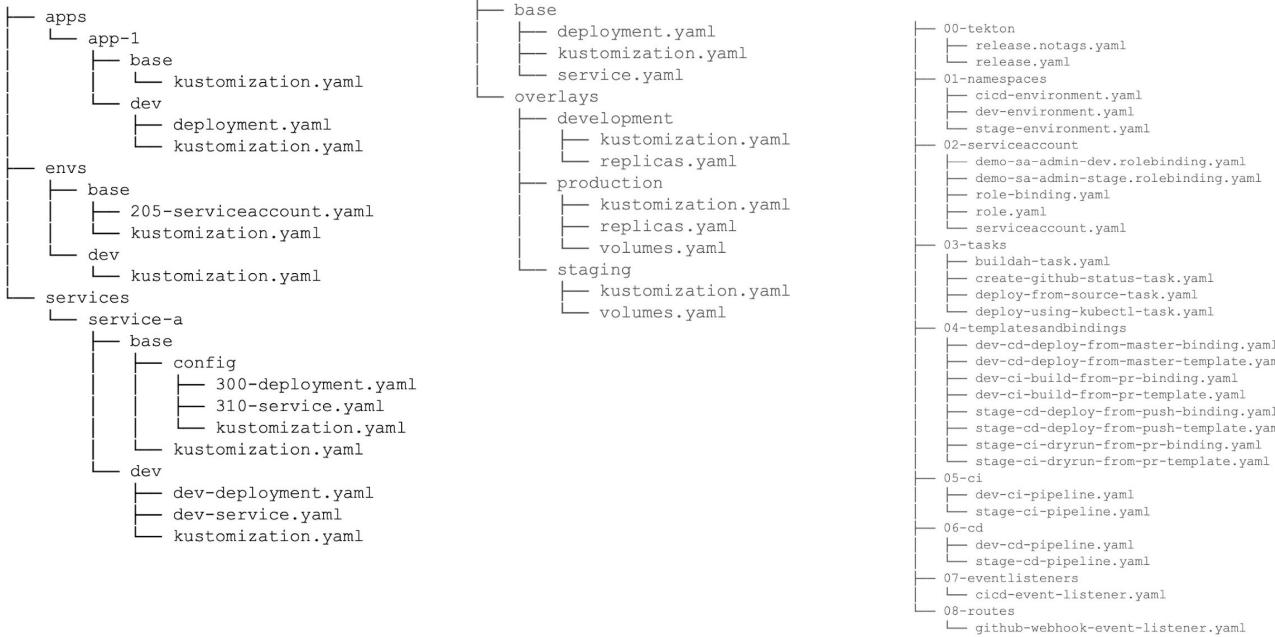
/taxi-config-dev.git

deploy

pipelines

README.md

Single Repository



ArgoCD - Order Dependent Deployments

- Sometimes you have cases where you need to deploy things in a specific order
 - Subscribe Operator before deploying instance
 - Create namespace before deploying app into it
 - Deploy required infrastructure before application
- Tools like kustomize and Helm might help handling this in some cases
- ArgoCD provides Sync Phases and Waves to address other use cases
 - 3 sync phases - Pre-sync, sync, post-sync
 - Each phase can have multiple waves, next wave does not proceed until previous phase is healthy

ArgoCD - Non-declarative requirements

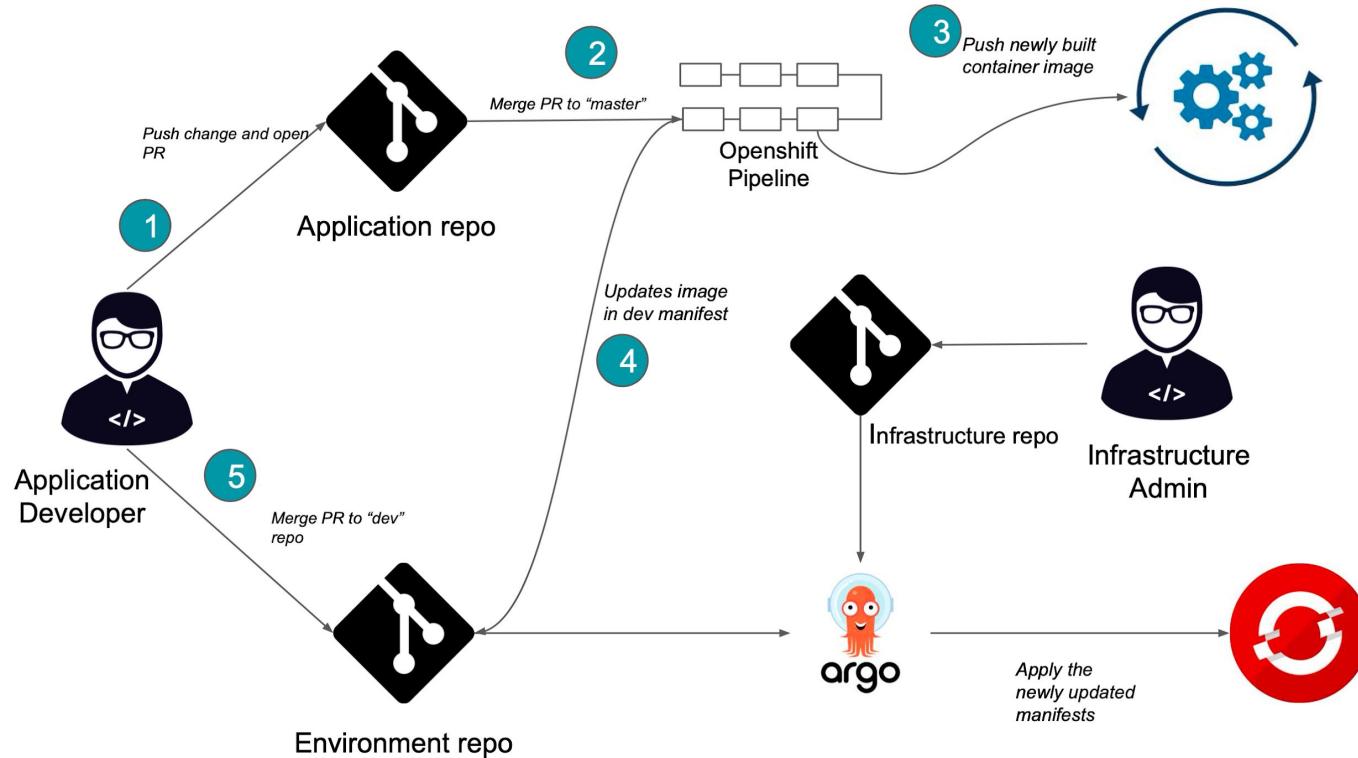
- There can be instances where you need to deploy something which cannot fully be done in a declarative way and it must be scripted
- Try minimizing this and leverage kubernetes primitives where possible
 - Init Containers
 - Jobs
 - Operators
- ArgoCD Resource Hooks
 - Hooks are ways to run scripts before, during and after a sync operation
 - Hooks can be run: PreSync, Sync, PostSync and SyncFail

ArgoCD - Integrating with CI/CD Tools

ArgoCD and NOT ArgoCI/CD

CI tools like Jenkins, Pipelines still required!

	ArgoCD Managed Deployment	Pipeline Managed Deployment
Pro	Consistent	Post-Test update of image reference
Con	Image reference updated in git before integration tests, manage rollback?	Inconsistent
Con	Pipeline tools must be able to wait for sync	



DEMO TIME

<https://github.com/wpernath/openshift-cicd-demo>

Summary

Summary

- You don't have to change anything, if you don't want
- You can use tools like Jenkins to do your CI
- OpenShift Builds is to integrate building your image in your existing pipeline
- OpenShift Pipelines (tekton) is a nice kubernetes-native way of pipelining
- GitOps and ArgoCD helps you to do a declarative approach
 - You as developer are used to use Git
 - You as admin are used to use Git
 - → Adoption of GitOps could be high



Thank you

 [linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)

 [facebook.com/redhatinc](https://www.facebook.com/redhatinc)

 [youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)

 twitter.com/RedHat